

OpenTelemetry Journey

2024.09

Jinwoong Kim

Who am I?

- Jinwoong Kim
- Cloud Architect @AWS Professional Services
- Speaker, Translator
- @ddiwoong

Observability Introduction

A *system* is **observable** if you can determine the *behavior* of the system based on its *outputs*.

Observability

Logs - Lines of text

```
hikari-pool-1 - Connection is not available, request timed out after 30000ms
```

Metrics - Time-ordered set of data

```
hikaricp_connections_timeout_total{pool="HikariPool-1",} 10.0
```

Tracing - Correlation analysis with Context

```
2022-05-28 18:09:04.165 INFO [service-b,757d0493f099b94b,4e8d66a6aa1c1ed6] 9989 ---
```

```
[nio-8686-exec-3] c.example.msaerrorresponse.BServiceApi : =====b-service=====
```

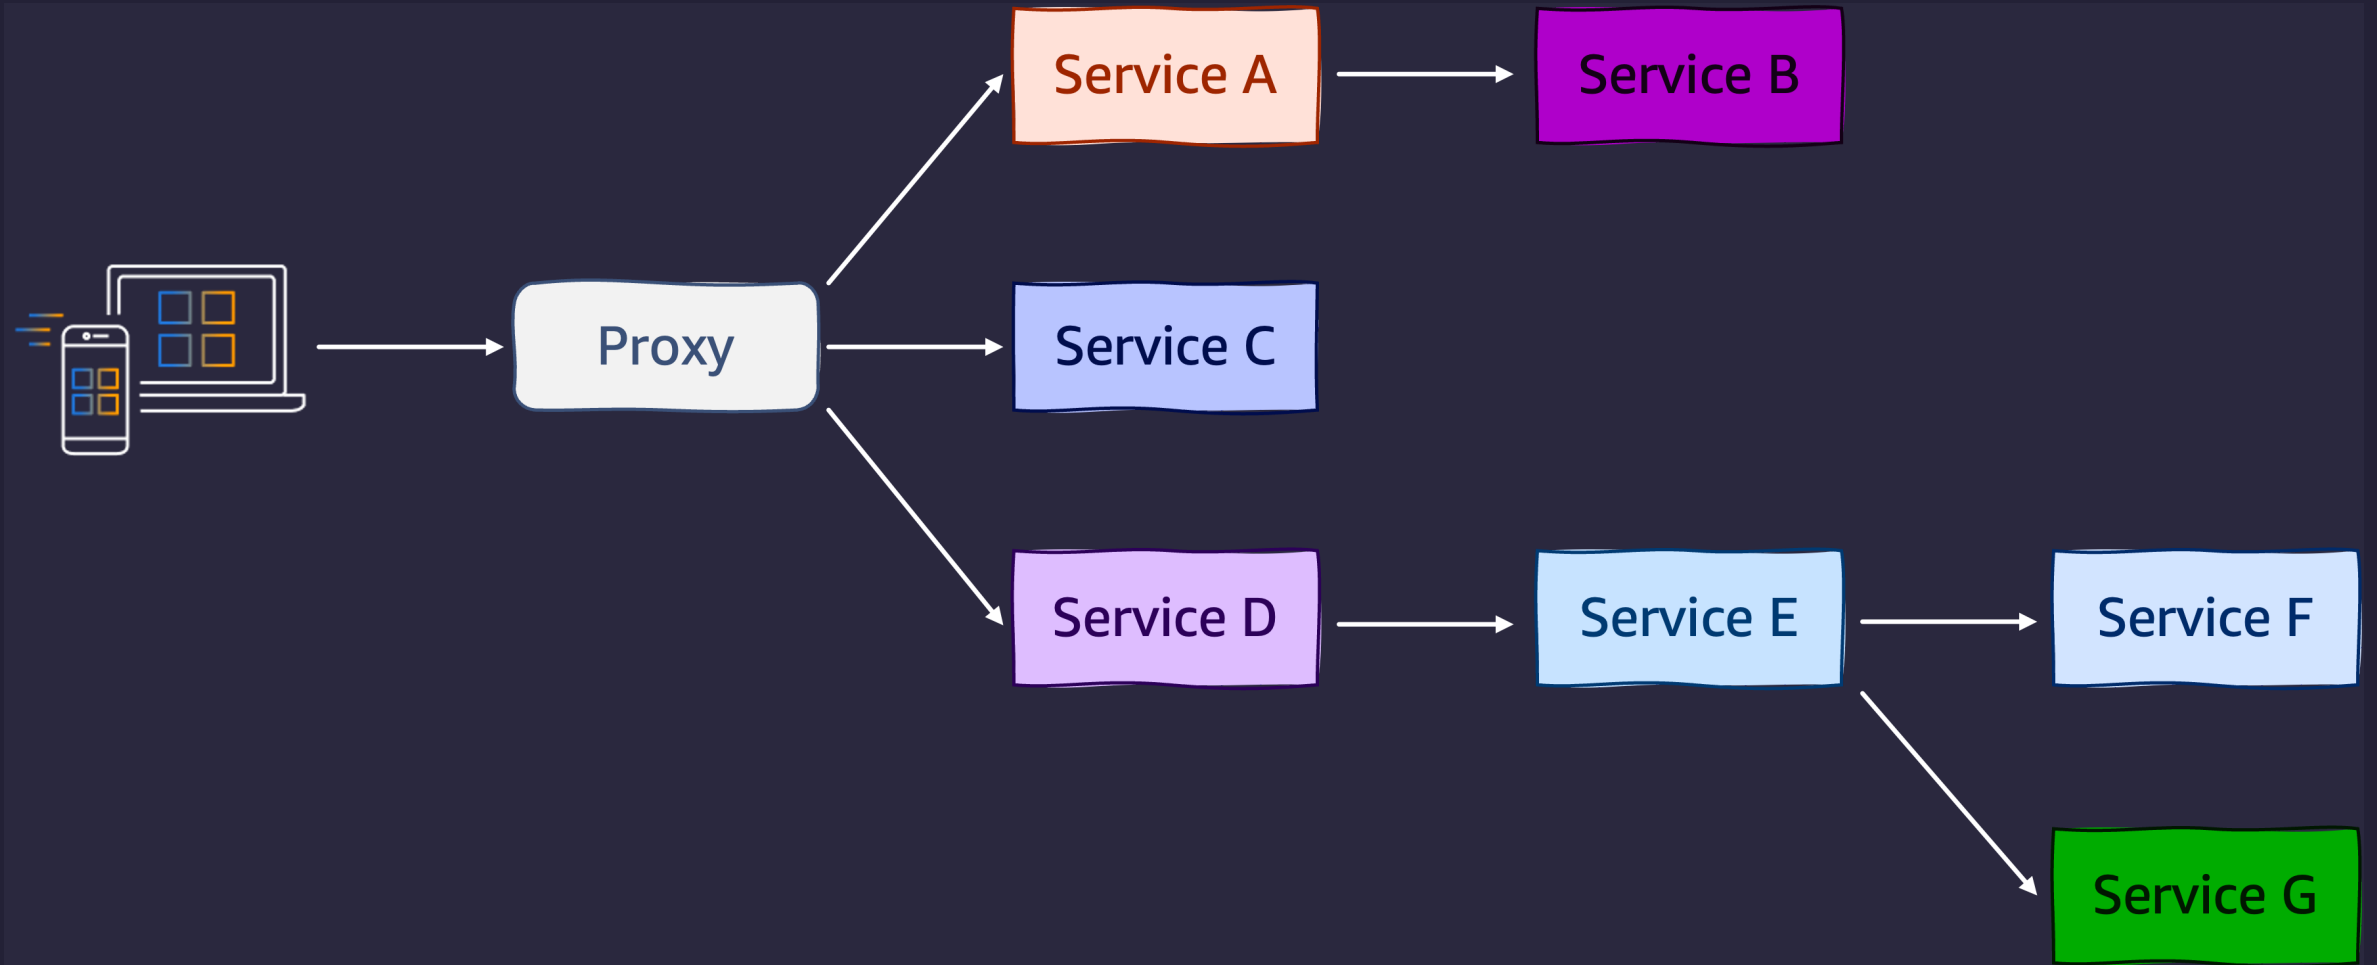
Distributed Trace

A **telemetry** method that indicates **latency** in specific parts of a system's process:

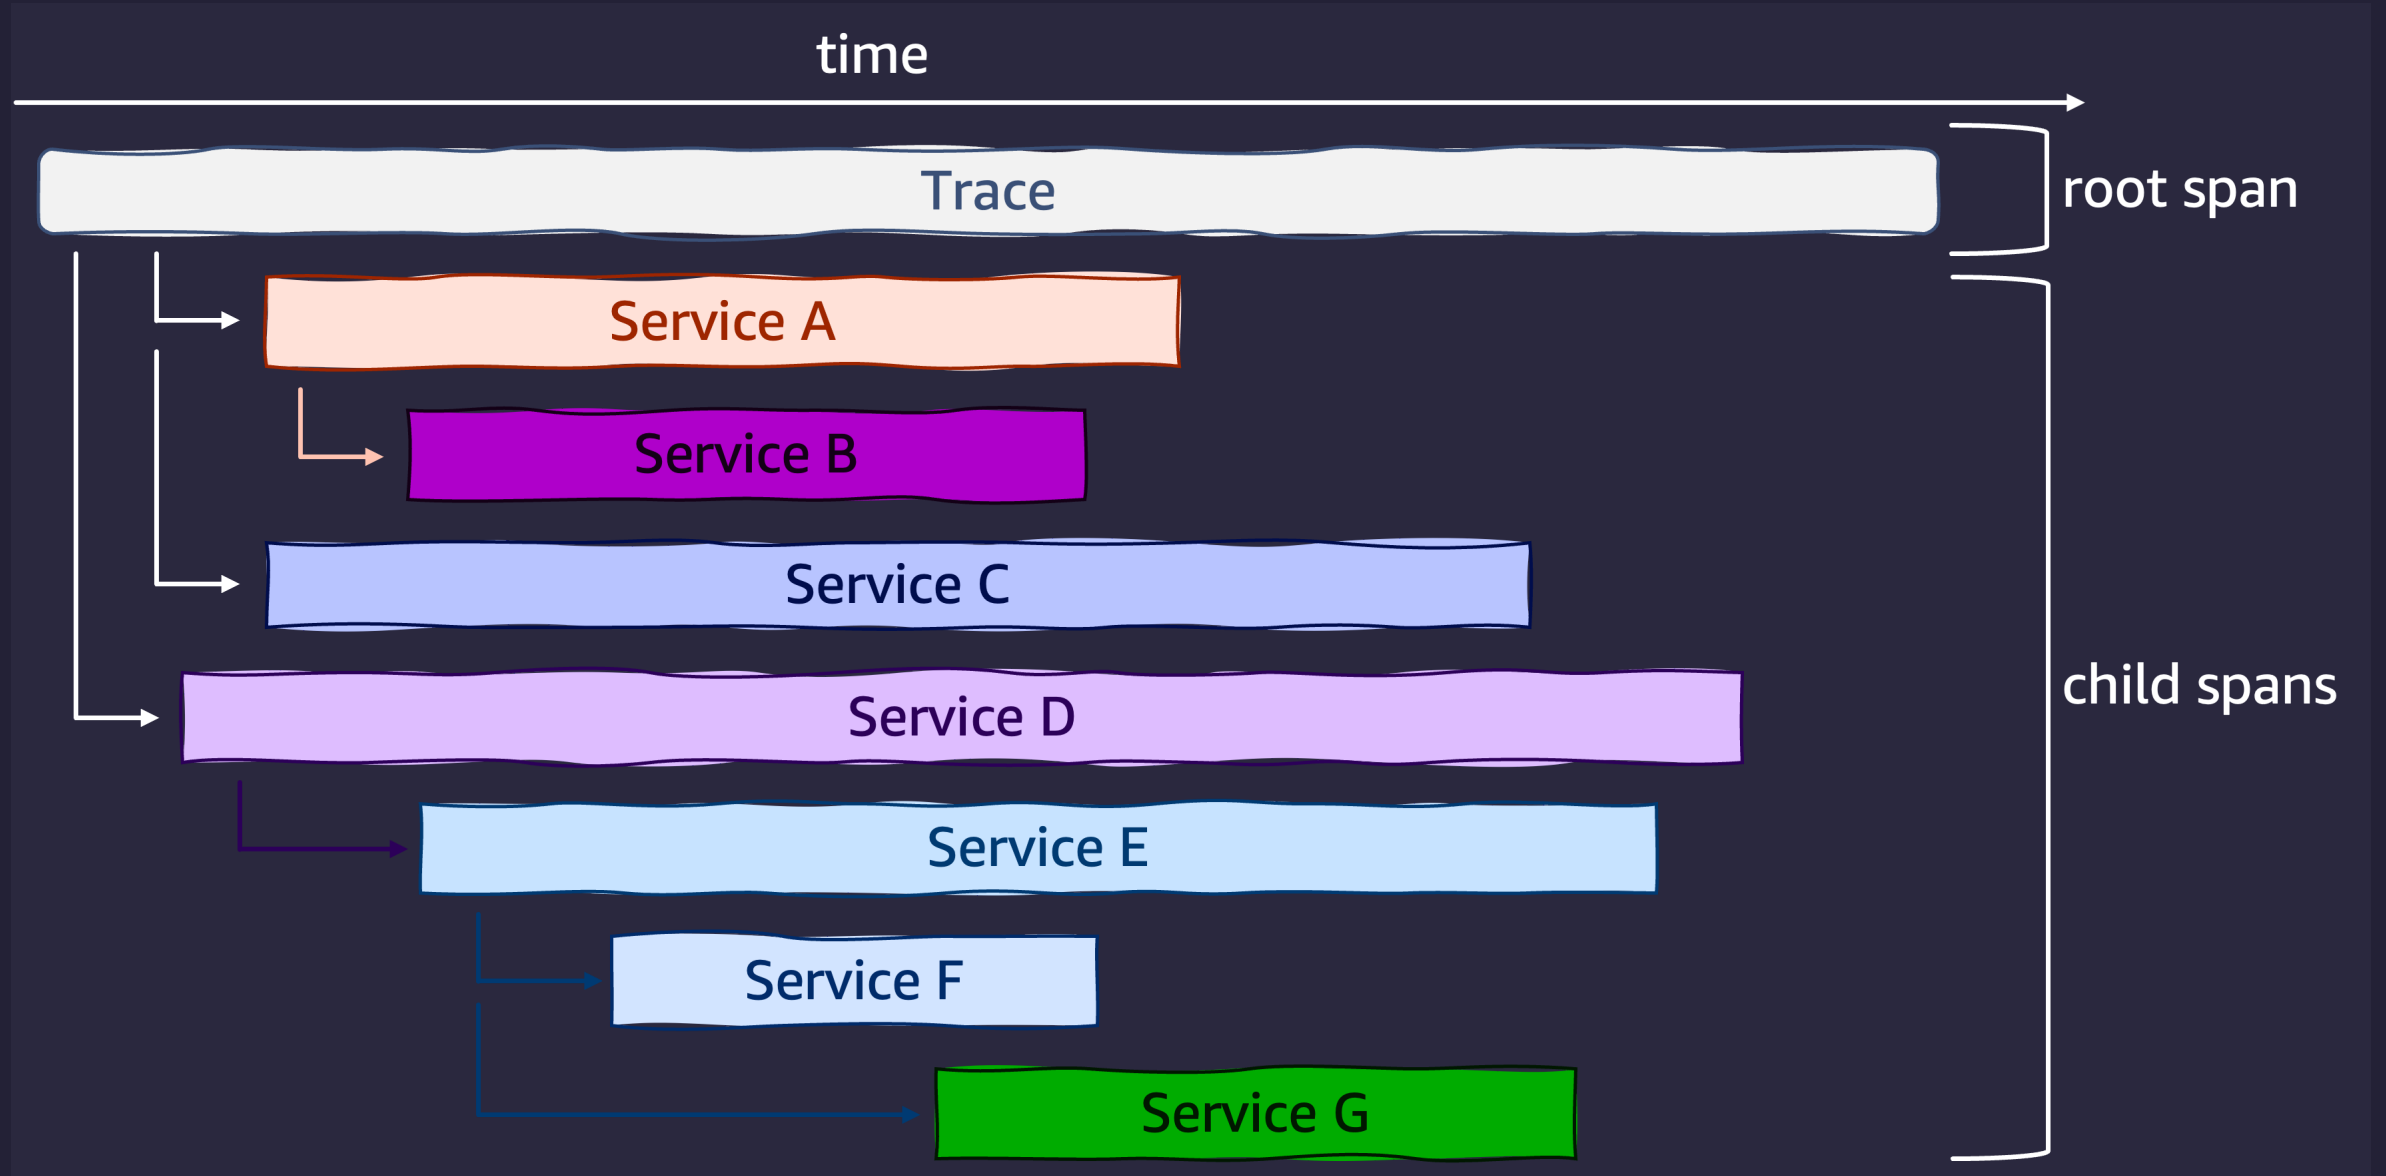
Records the path that requests take as they propagate through microservices and serverless architectures.

It is crucial for **observability** because it measures **dependencies** and **relationships** among numerous components in modern architectures like microservices environments, identifying **latency bottlenecks**.

Trace

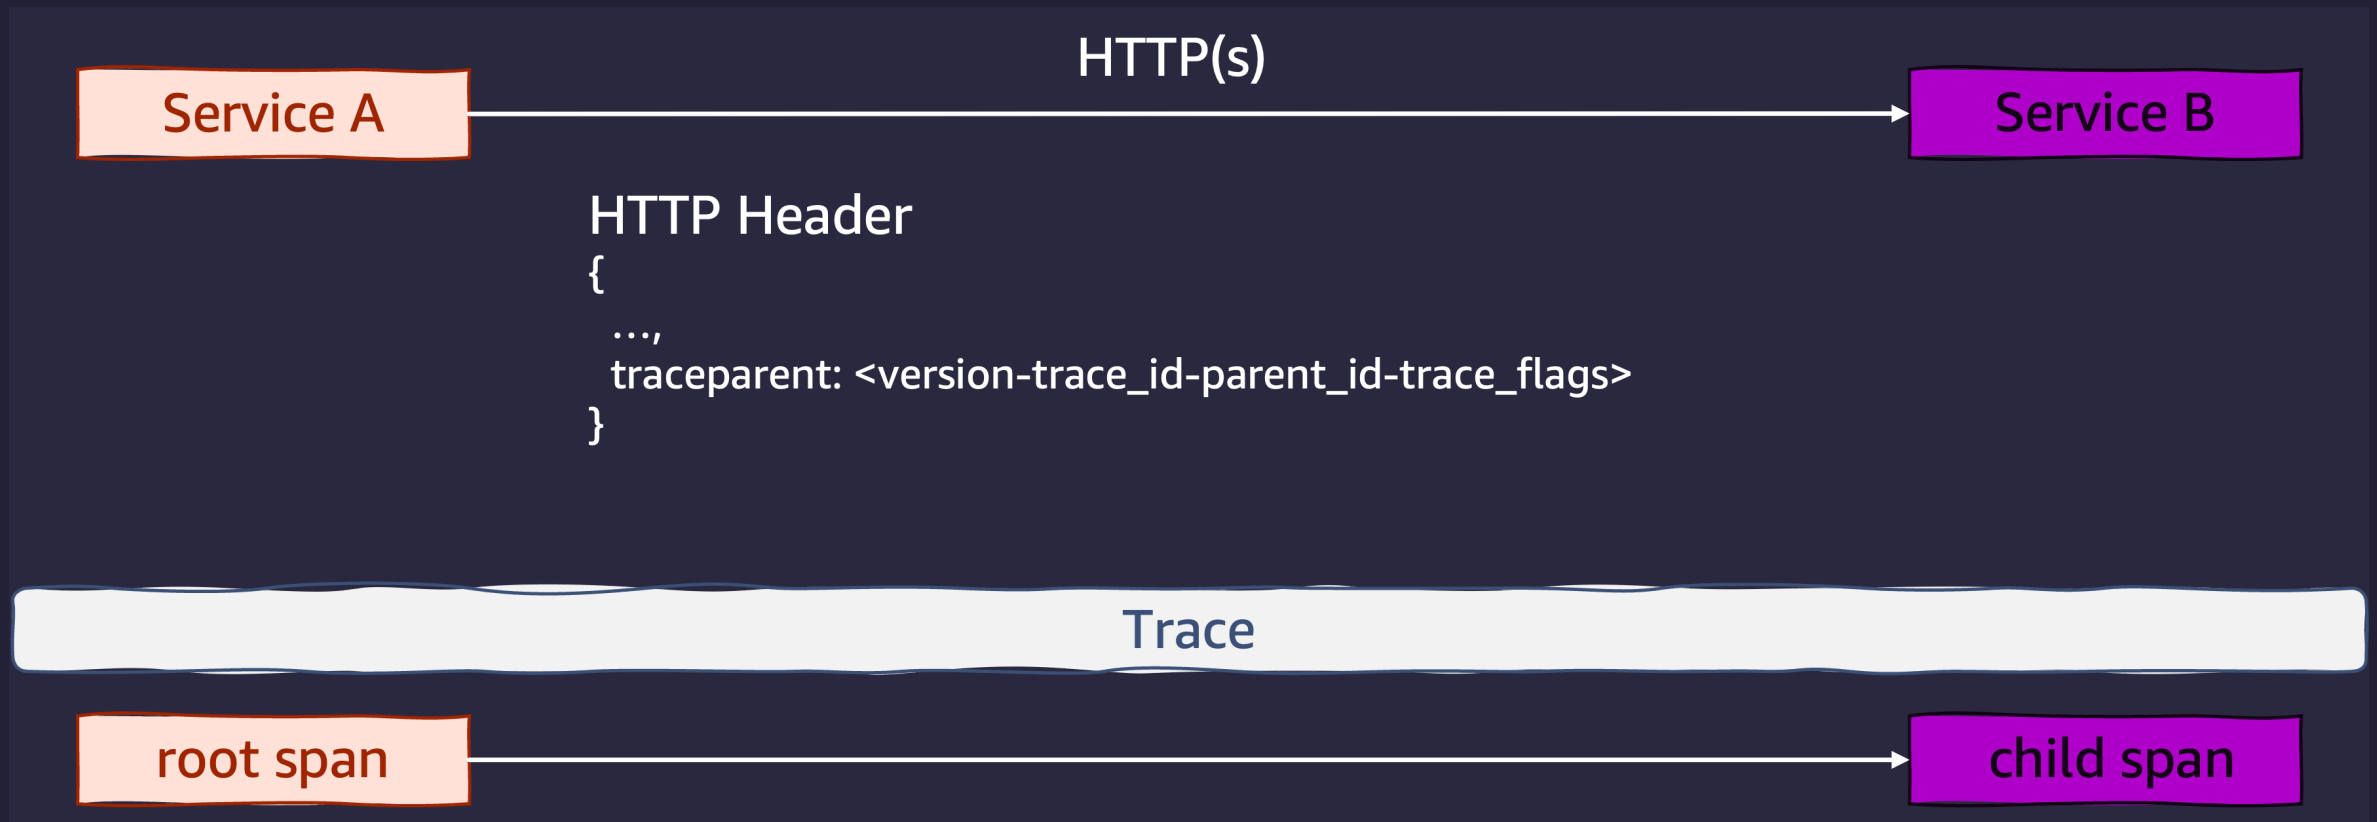


Trace (Span)

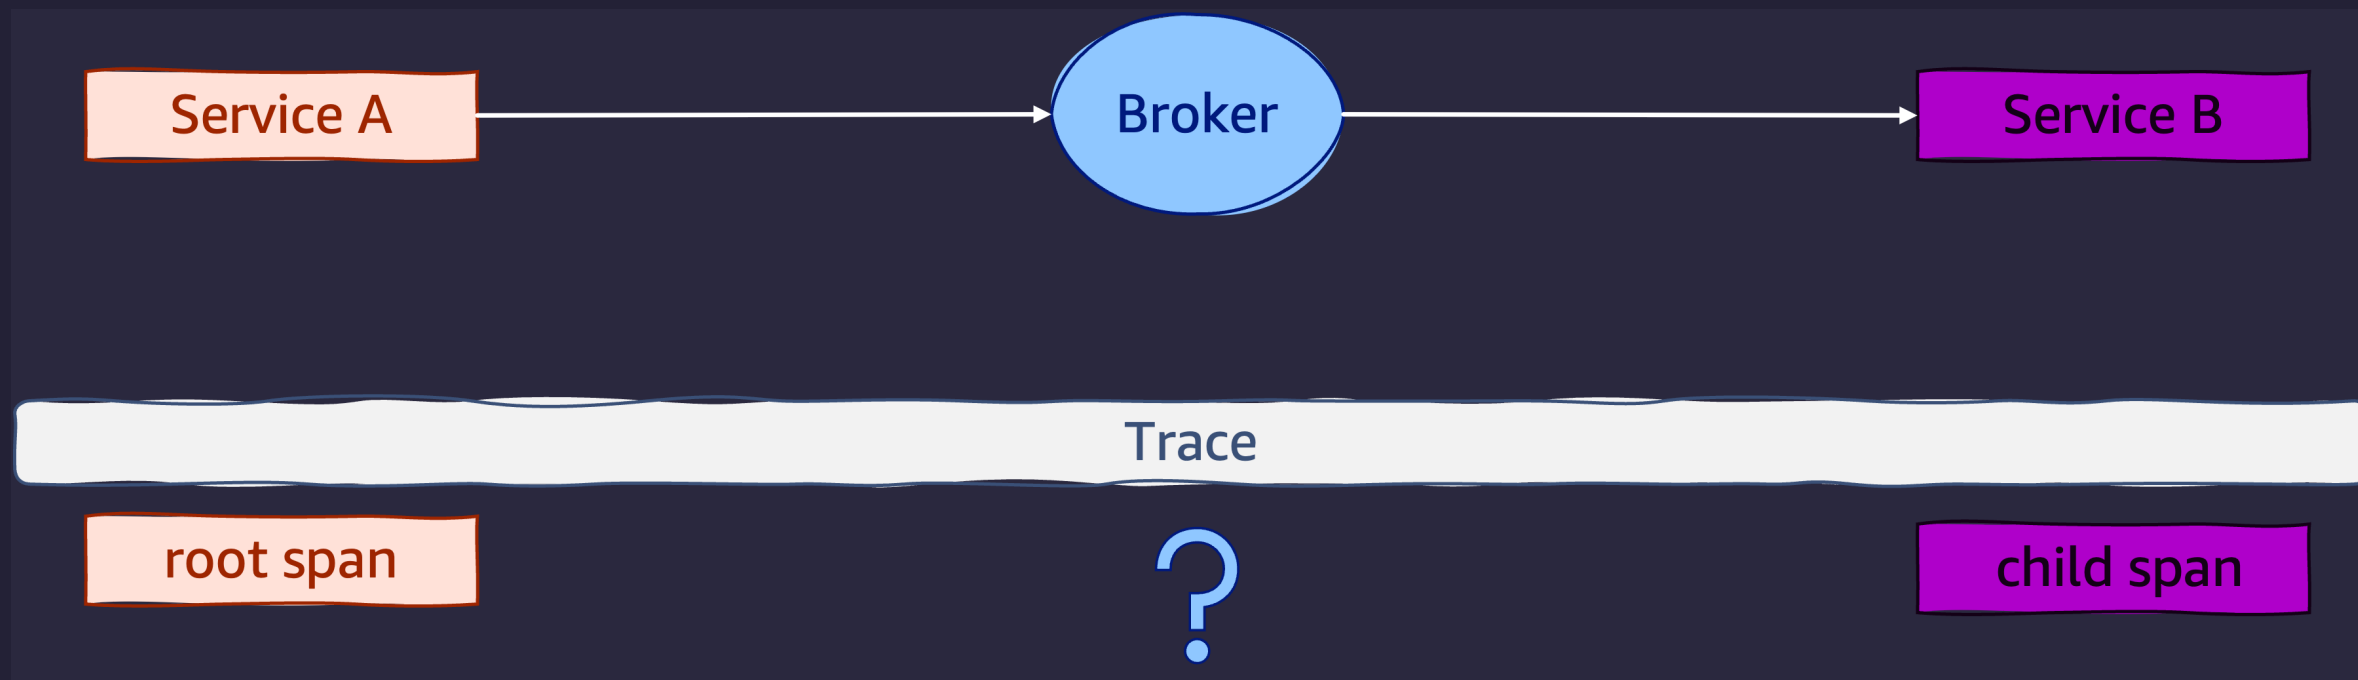


Trace - context propagation

W3C Trace Context : <https://w3c.github.io/trace-context/>



Trace - context propagation



Trace - context propagation



Trace

root span

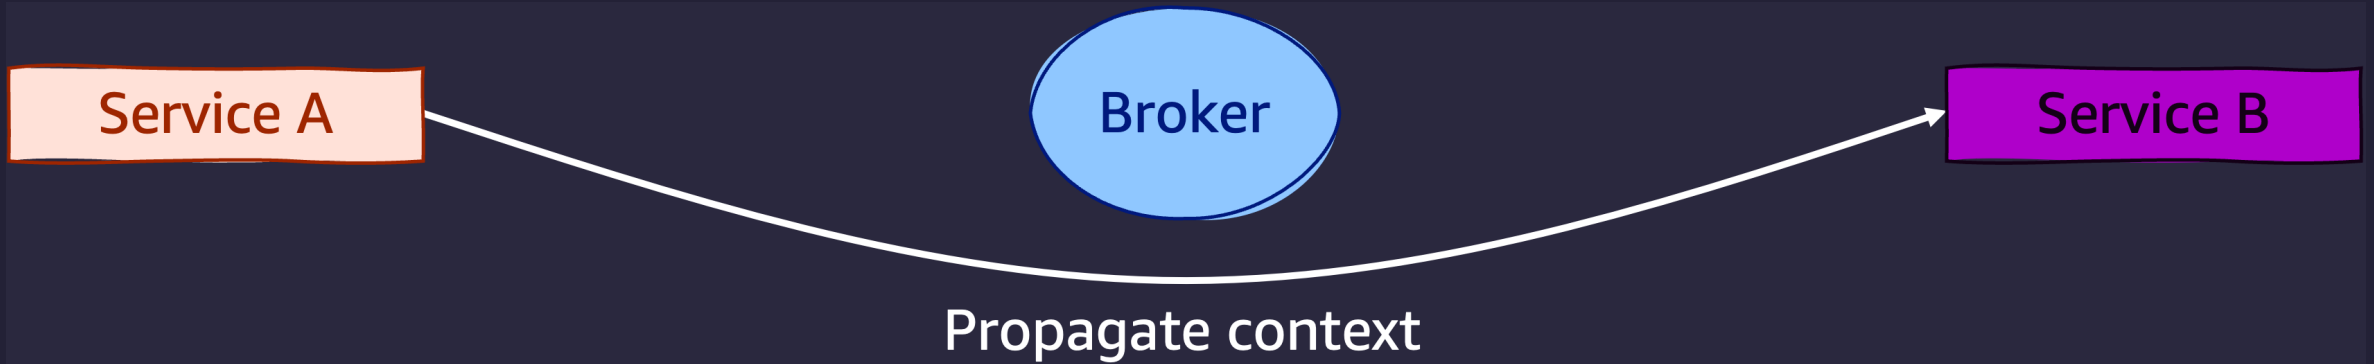
```
{
  "name": "root-span",
  "context": {
    "trace_id": "0x5b8aa5a2d2c872e8321cf37308d69df2",
    "span_id": "0x051581bf3cb55c13"
  }
  "parent_id": null,
  "start_time": "2022-04-29T18:52:58.114201Z",
  "end_time": "2022-04-29T18:52:58.114687Z"
}
```

?

child span

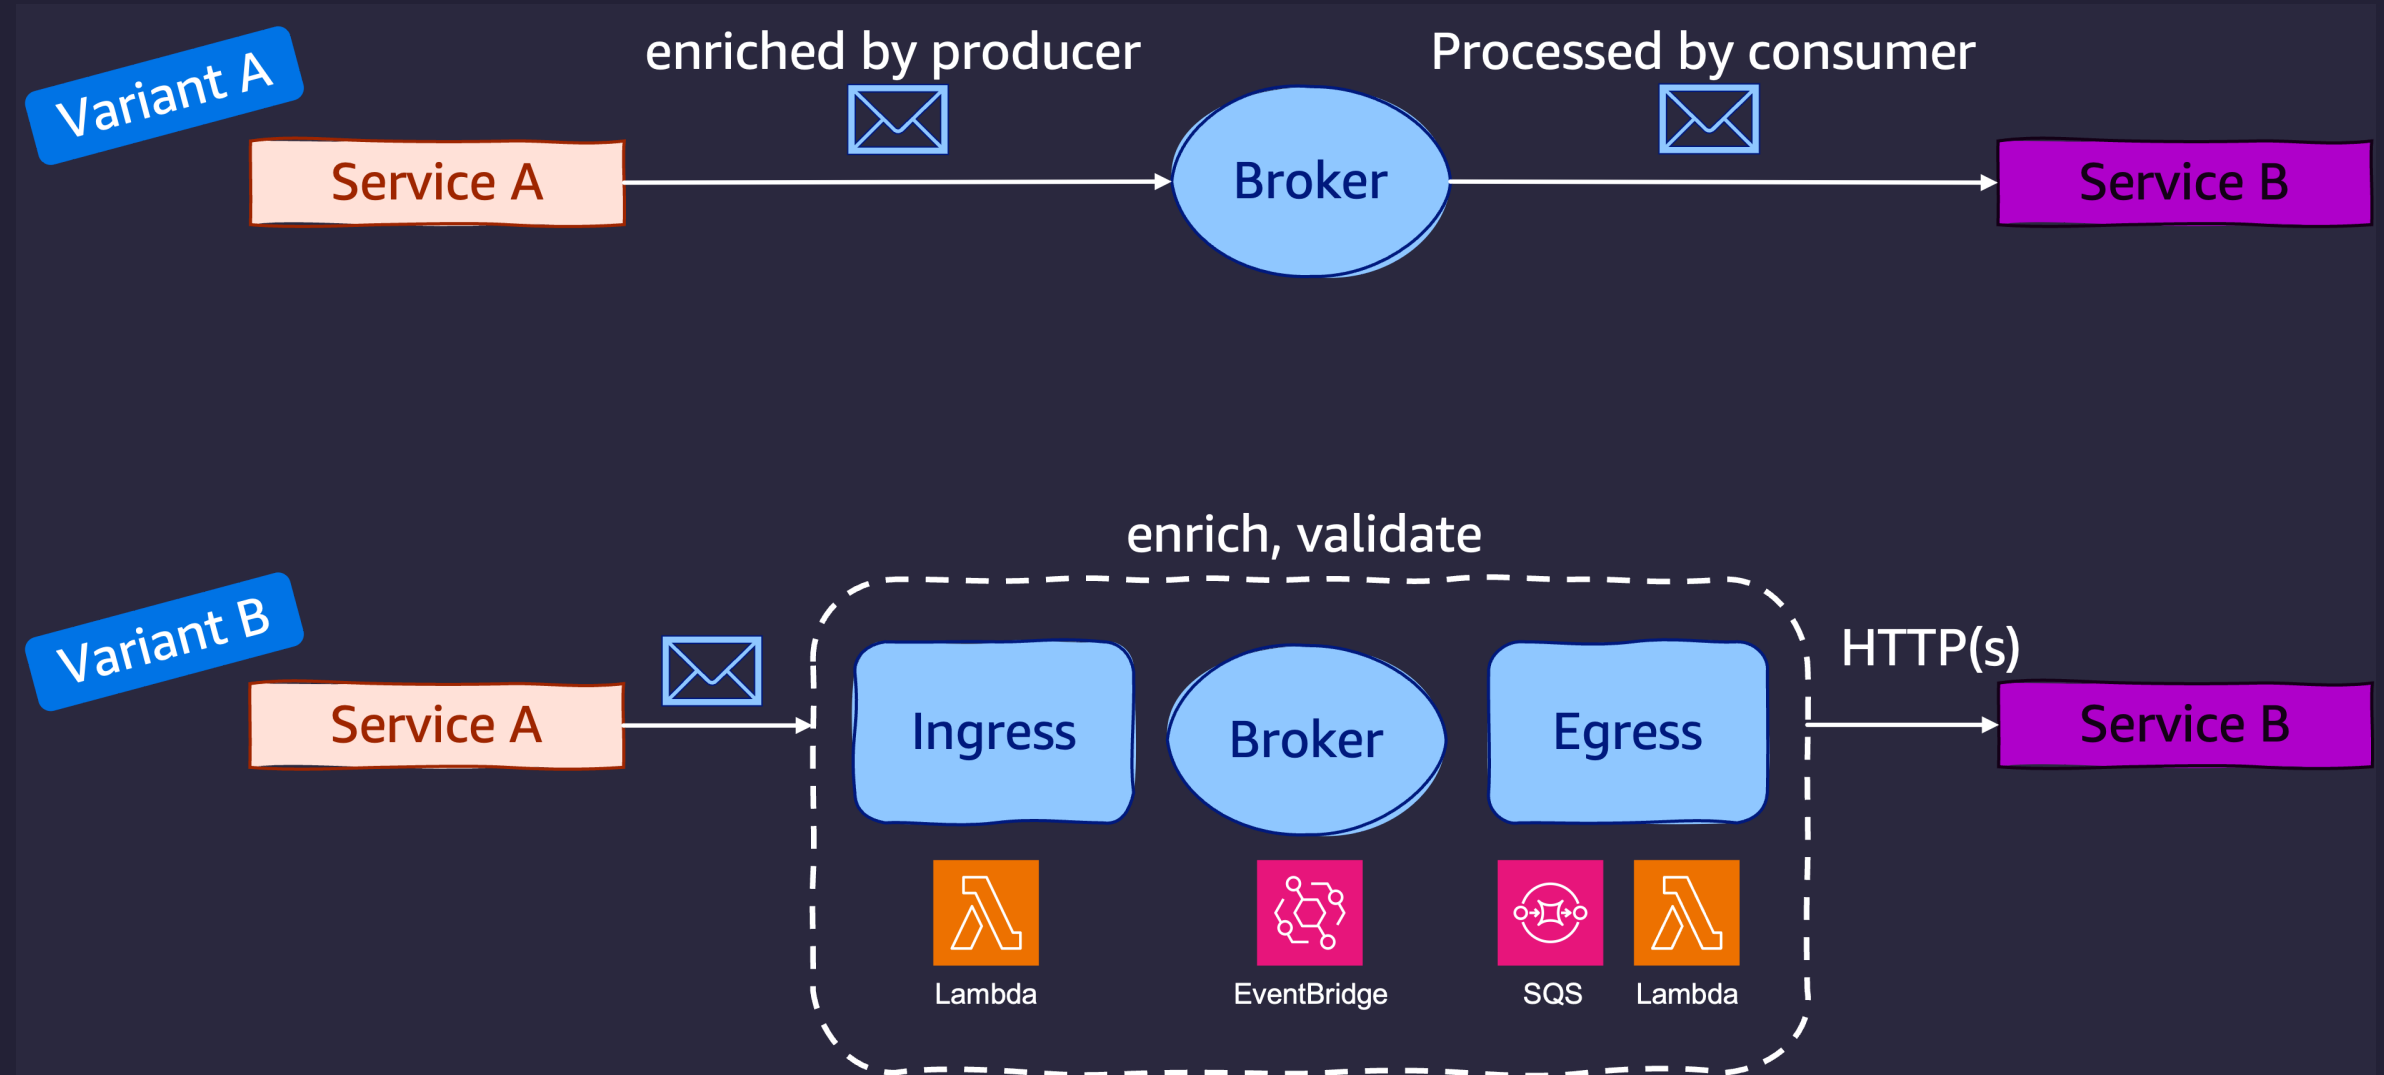
```
{
  "name": "child-span",
  "context": {
    "trace_id": "0x5b8aa5a2d2c872e8321cf37308d69df2",
    "span_id": "0x5fb397be34d26b51"
  }
  "parent_id": "0x051581bf3cb55c13",
  "start_time": "2022-04-29T18:52:58.114504Z",
  "end_time": "2022-04-29T22:52:58.114561Z"
}
```

Trace - context propagation



```
{  
  → "metadata": {  
    < some mandatory fields like timestamps, domain information, ...>  
    "otel": {  
      "trace_parent": {  
        "trace_id": "0x5b8aa...",  
        "span_id": "0x05158..."  
      }  
    }  
  },  
  → "data": {  
    < original producer event >  
  }  
}
```

Trace - context propagation



Span

```
{
  "name": "/v1/sys/health",
  "context": {
    "trace_id": "7bba9f33312b3dbb8b2c2c62bb7abe2d",
    "span_id": "086e83747d0e381e"
  },
  "parent_id": "",
  "start_time": "2021-10-22 16:04:01.209458162 +0000 UTC",
  "end_time": "2021-10-22 16:04:01.209514132 +0000 UTC",
  "status_code": "STATUS_CODE_OK",
  "status_message": "",
  "attributes": {
    "http.scheme": "http",
    "http.host": "10.177.2.152:26040",
  },
  "events": [
    {
      "name": "",
      "message": "OK",
      "timestamp": "2021-10-22 16:04:01.209512872 +0000 UTC"
    }
  ]
}
```

Span

170 / 521 requests | 2.5 MB / 2.8 MB transferred | 32.8 MB / 34.9 MB resources | Finish: 18.72 s | DOMContentLoaded: 1.66 s | Load: 2.54 s

```
{
  "attributes": {
    "response_size": "2.5",
    "content_size": "2.8",
  },
  "events": [
    {
      "name": "Finish",
      "message": "OK",
      "timestamp": "2021-10-22 16:04:01.209512872 +0000 UTC"
    }
  ]
}
```

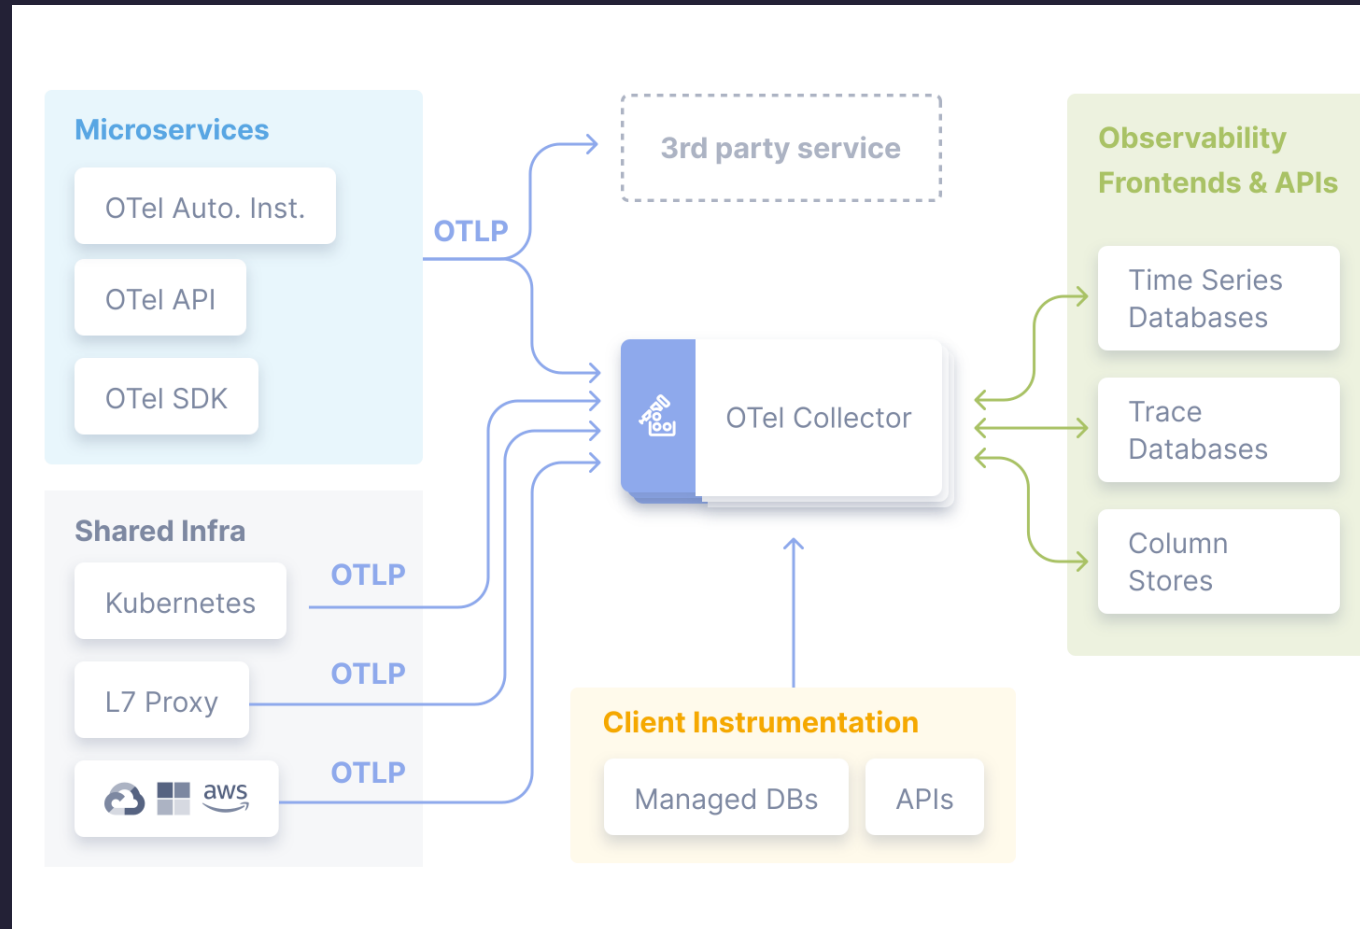


<https://opentelemetry.io/docs/specs/otel/trace/exceptions/>

Hello, OpenTelemetry

Open source project hosted on CNCF Specifications, Implementations for instrumentation and transmissions of telemetry data (metrics, logs, traces)

1. Cross-language specifications
2. OpenTelemetry Collector (agent)
3. SDKs for each language
4. Auto Instrumentation



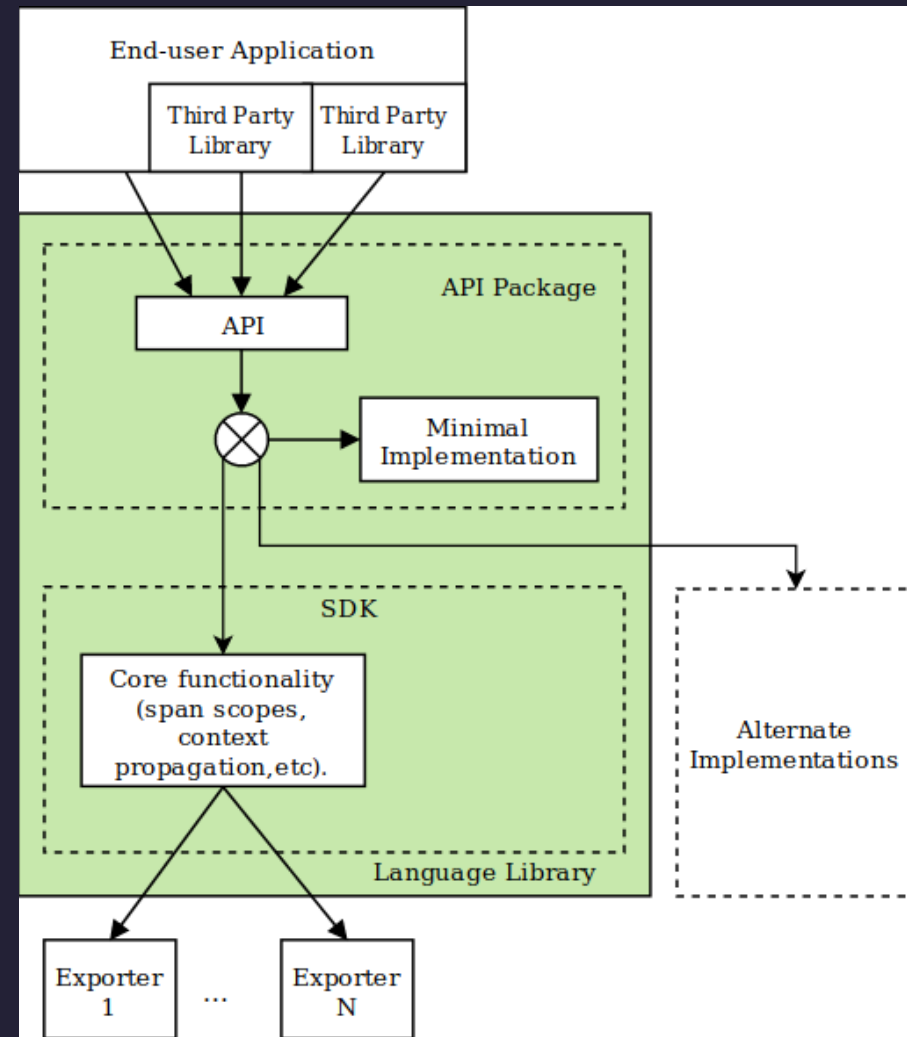
OpenTelemetry Instrumentation

1. Code-based solutions via official APIs and SDKs for most languages

- **API** defines data types and how to generate telemetry data.
- **SDK** defines a language-specific implementation of the API, plus configuration, data processing and exporting.

2. Zero-code solutions

- Go, .NET, PHP, Python, Java, JavaScript



Manually Instrumentation (Python)

```
@app.route("/server_request")
def server_request():
    with tracer.start_as_current_span(
        "server_request",
        context=extract(request.headers),
        kind=trace.SpanKind.SERVER,
        attributes=collect_request_attributes(request.environ),
    ):
        print(request.args.get("param"))
        return "served"
```

Programmatically-instrumented server (Python)

```
instrumentor = FlaskInstrumentor()

app = Flask(__name__)

instrumentor.instrument_app(app)
# instrumentor.instrument_app(app, excluded_urls="/server_request")
@app.route("/server_request")
def server_request():
    print(request.args.get("param"))
    return "served"
```

OpenTelemetry Registry

OpenTelemetry instrumentation libraries

<https://opentelemetry.io/ecosystem/registry/>

Search **NGINX**

The screenshot shows the OpenTelemetry Registry website. The header includes the OpenTelemetry logo and navigation links: Docs, Ecosystem, Status, Community, Blog, English, and a search bar. The main heading is "Registry" with the subtitle "Find libraries, plugins, integrations, and other useful tools for using and extending OpenTelemetry." Below this, a search bar shows "Search 838 entries" for the term "nginx". A description states: "The OpenTelemetry Registry allows you to search for instrumentation libraries, collector components, utilities, and other useful projects in the OpenTelemetry ecosystem. If you are a project maintainer, you can [add your project to the OpenTelemetry Registry](#)." The search results list two items:

- NGINX Instrumentation** by [OpenTelemetry Authors](#). Description: "NGINX OpenTelemetry module to add OpenTelemetry distributed tracing support to NGINX." Metadata: C++ Language, Instrumentation Component, Apache 2.0 License. Link: [Repository](#).
- NGINX Native OTel Module** by [NGINX, Inc.](#). Metadata: ★ new ♥ first party integration, C++ Language, Instrumentation Component, Apache-2.0 License. Links: [Website](#), [Repository](#).

OpenTelemetry collector



OTEL COLLECTOR

OpenTelemetry protocol (OTLP)

<https://github.com/open-telemetry/opentelemetry-proto/blob/main/docs/specification.md>

OTLP is implemented over **gRPC** and **HTTP** transports and specifies the Protocol Buffers schema used for payloads.

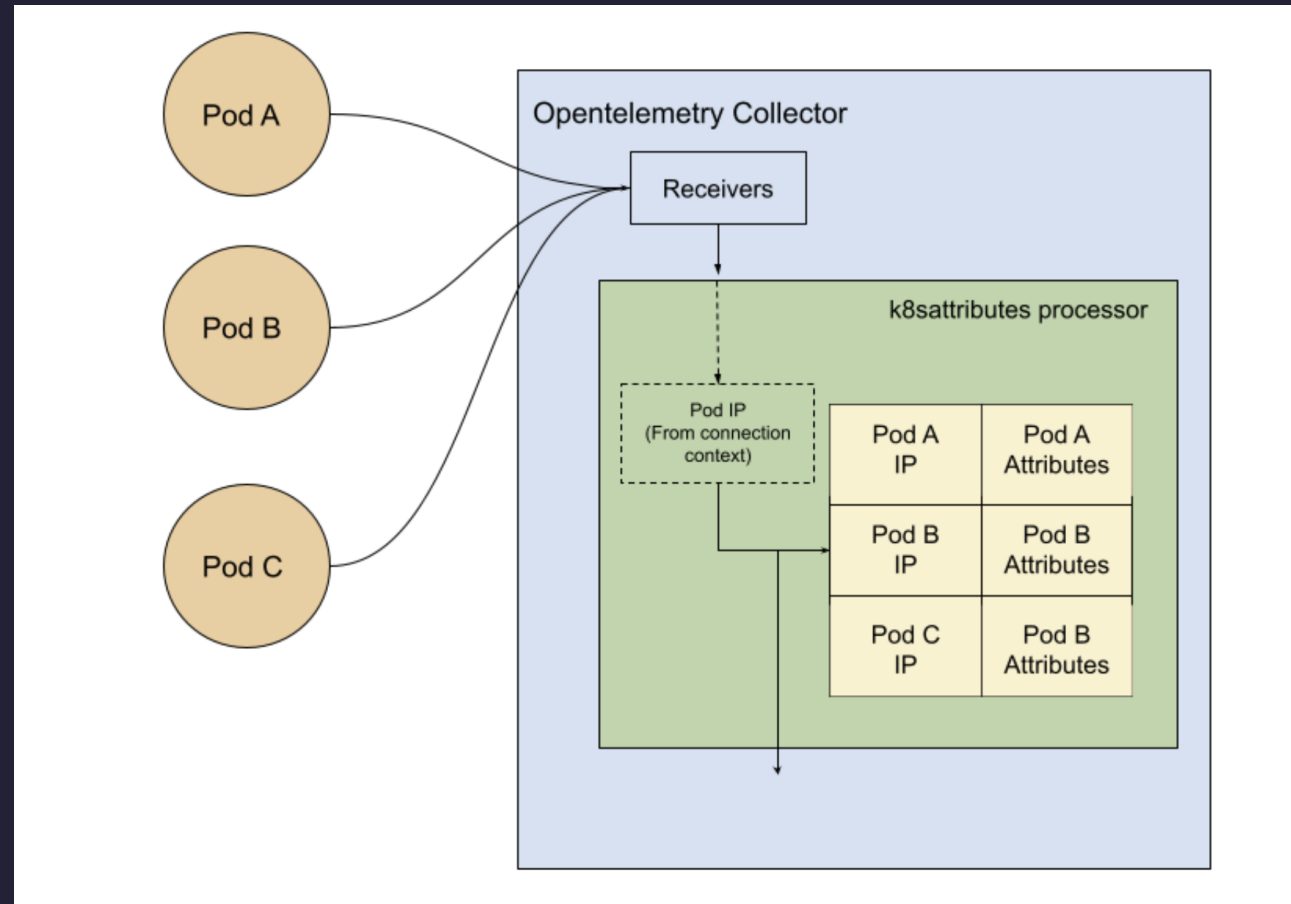
OTLP is a request/response style protocol where **clients** send requests and the **server** replies with corresponding responses.

All server components must support the following transport compression options: **none**, **gzip**

Resource Semantic Conventions

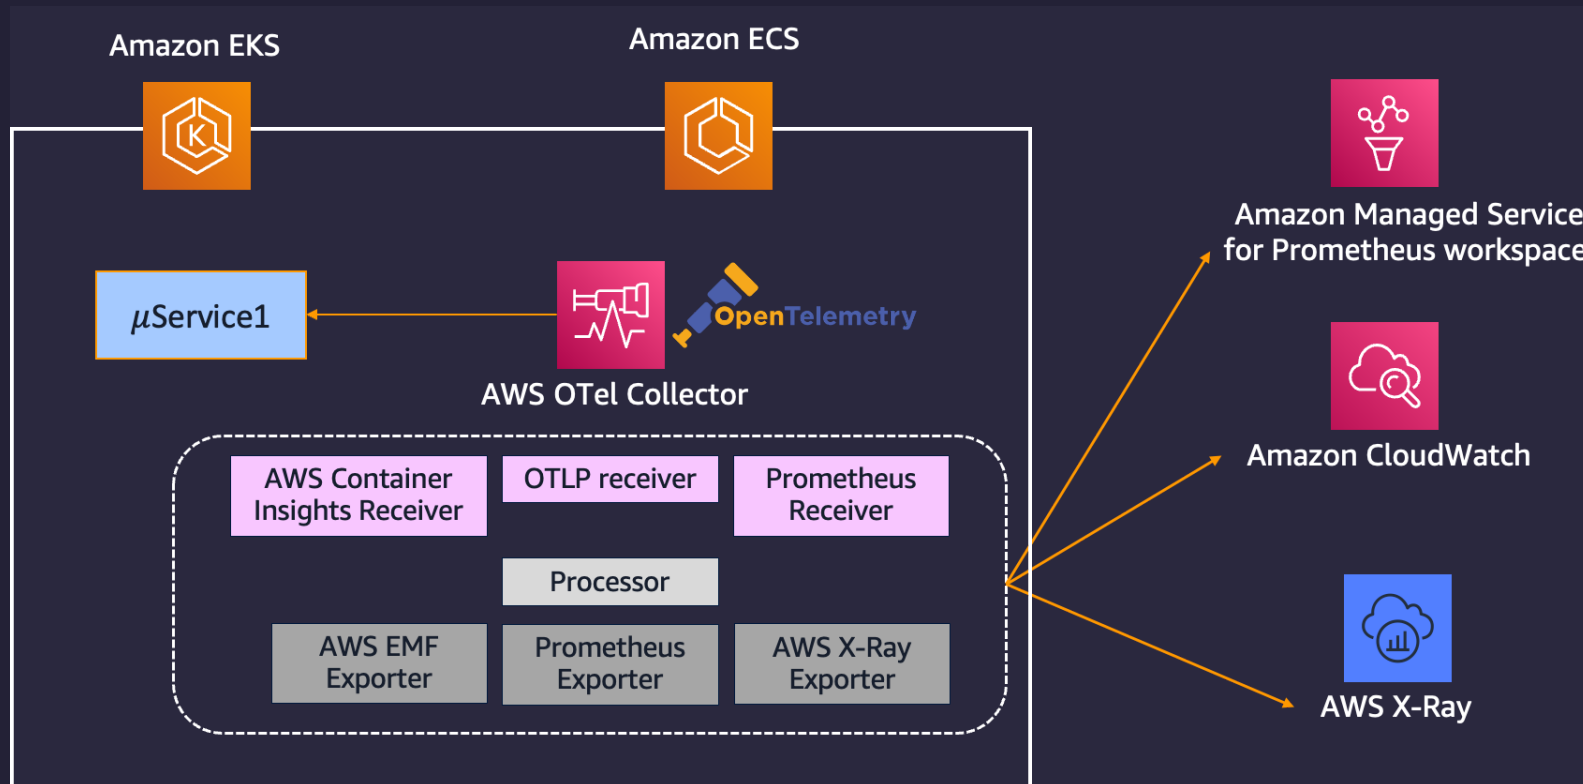
A **Resource** represents the entity producing telemetry as resource attributes.

<https://opentelemetry.io/docs/specs/semconv/>

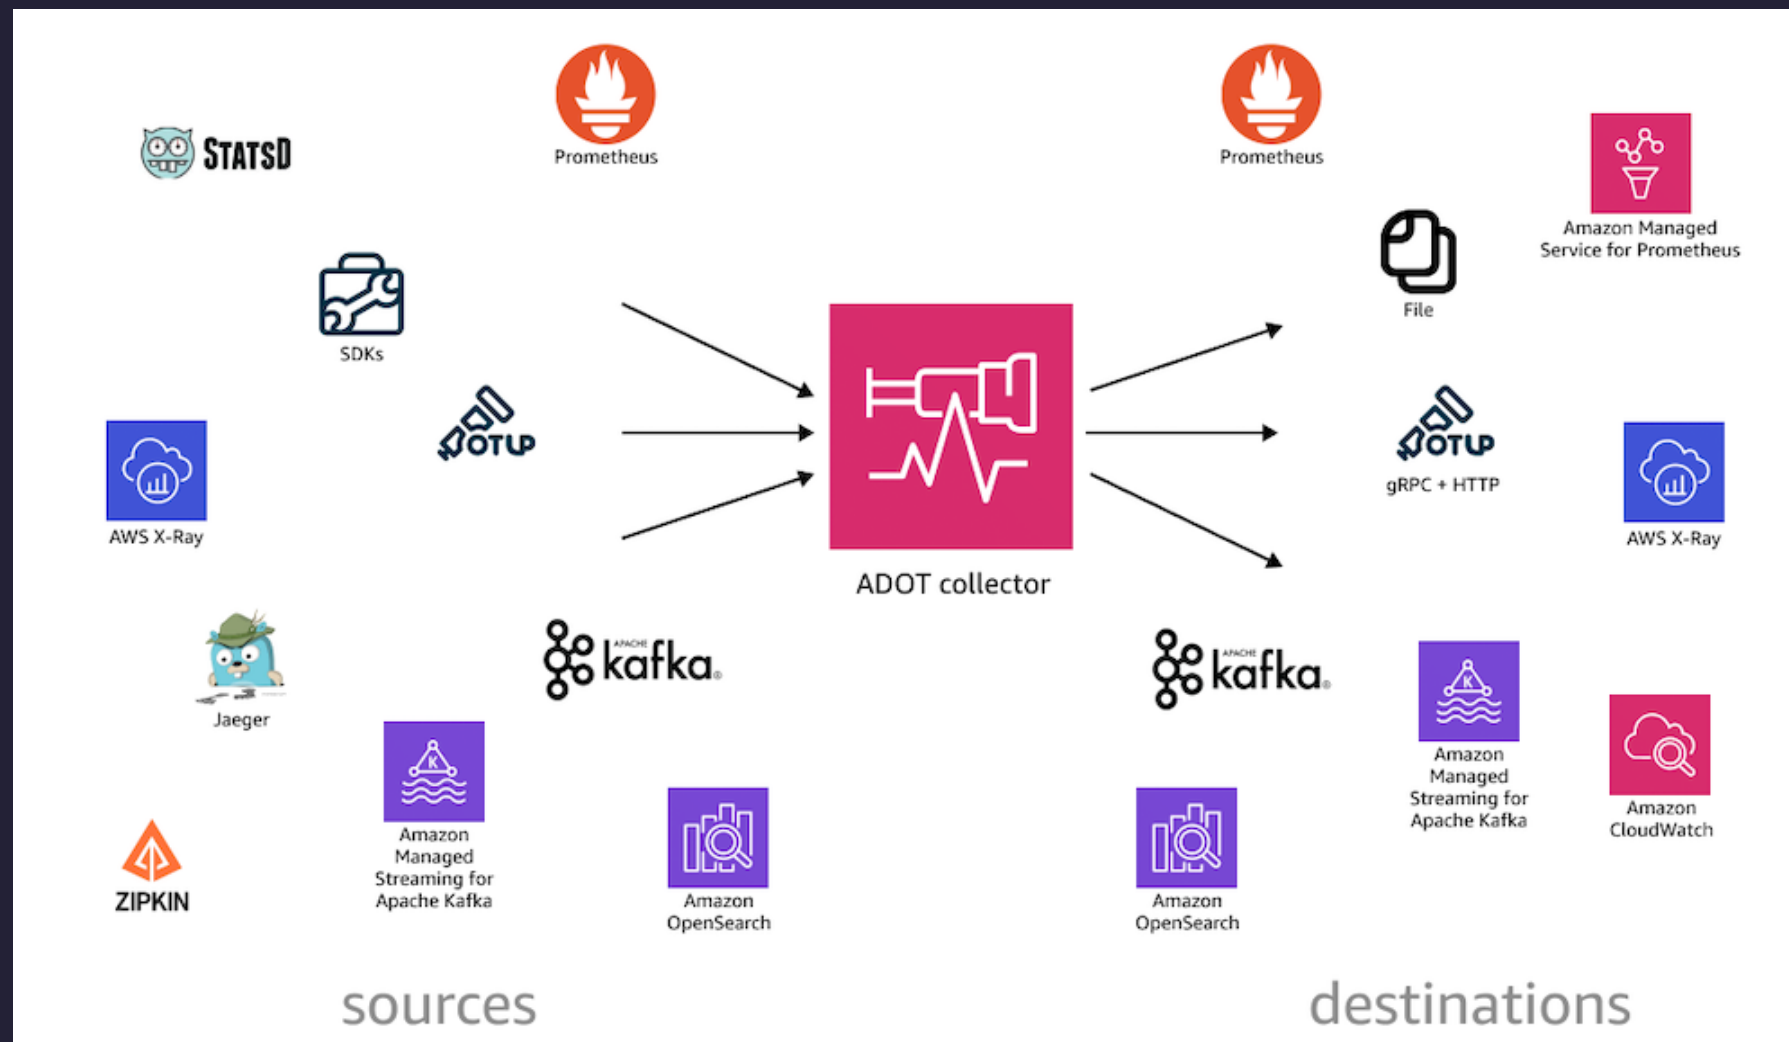


AWS Distro for OpenTelemetry (ADOT)

- Secure, production ready, and supported by AWS OpenTelemetry distribution
- Extend functionality for ease of use on AWS




AWS Distro for OpenTelemetry (ADOT)



OpenTelemetry Demo

OpenTelemetry Demo is composed of microservices written in different programming languages that talk to each other over gRPC and HTTP; and a load generator which uses Locust to fake user traffic.

- Web store
- Grafana
- Load Generator
- Jaeger UI



New Otel Feature

Envoy and Istio

Profiling Agent

LLM Observability

OpenTelemetry Collector Antipatterns

References

- <https://opentelemetry.io/docs/>
- <https://w3c.github.io/trace-context/>
- <https://w3c.github.io/baggage/>
- <https://github.com/open-telemetry/opentelemetry-specification>
- <https://opentelemetry.io/docs/specs/semconv/>
- <https://opentelemetry.io/docs/specs/otel/protocol/>
- <https://opentelemetry.io/docs/concepts/sampling/>
- <https://opentelemetry.io/docs/demo/>
- <https://opentelemetry.io/blog/2024/>

Thank You