

OpenTelemetry Journey

2024.09

Jinwoong Kim

Who am I?

- Jinwoong Kim
- Cloud Architect @AWS Professional Services
- Speaker, Translator
- @ddiwoong

Observability Introduction

A *system* is **observable** if you can determine the *behavior* of the system based on its *outputs*.

Observability

Logs - Lines of text

```
hikari-pool-1 - Connection is not available, request timed out after 30000ms
```

Metrics - Time-ordered set of data

```
hikaricp_connections_timeout_total{pool="HikariPool-1",} 10.0
```

Tracing - Correlation analysis with Context

```
2022-05-28 18:09:04.165 INFO [service-b,757d0493f099b94b,4e8d66a6aa1c1ed6] 9989 ---
```

```
[nio-8686-exec-3] c.example.msaerrorresponse.BServiceApi : =====b-service=====
```

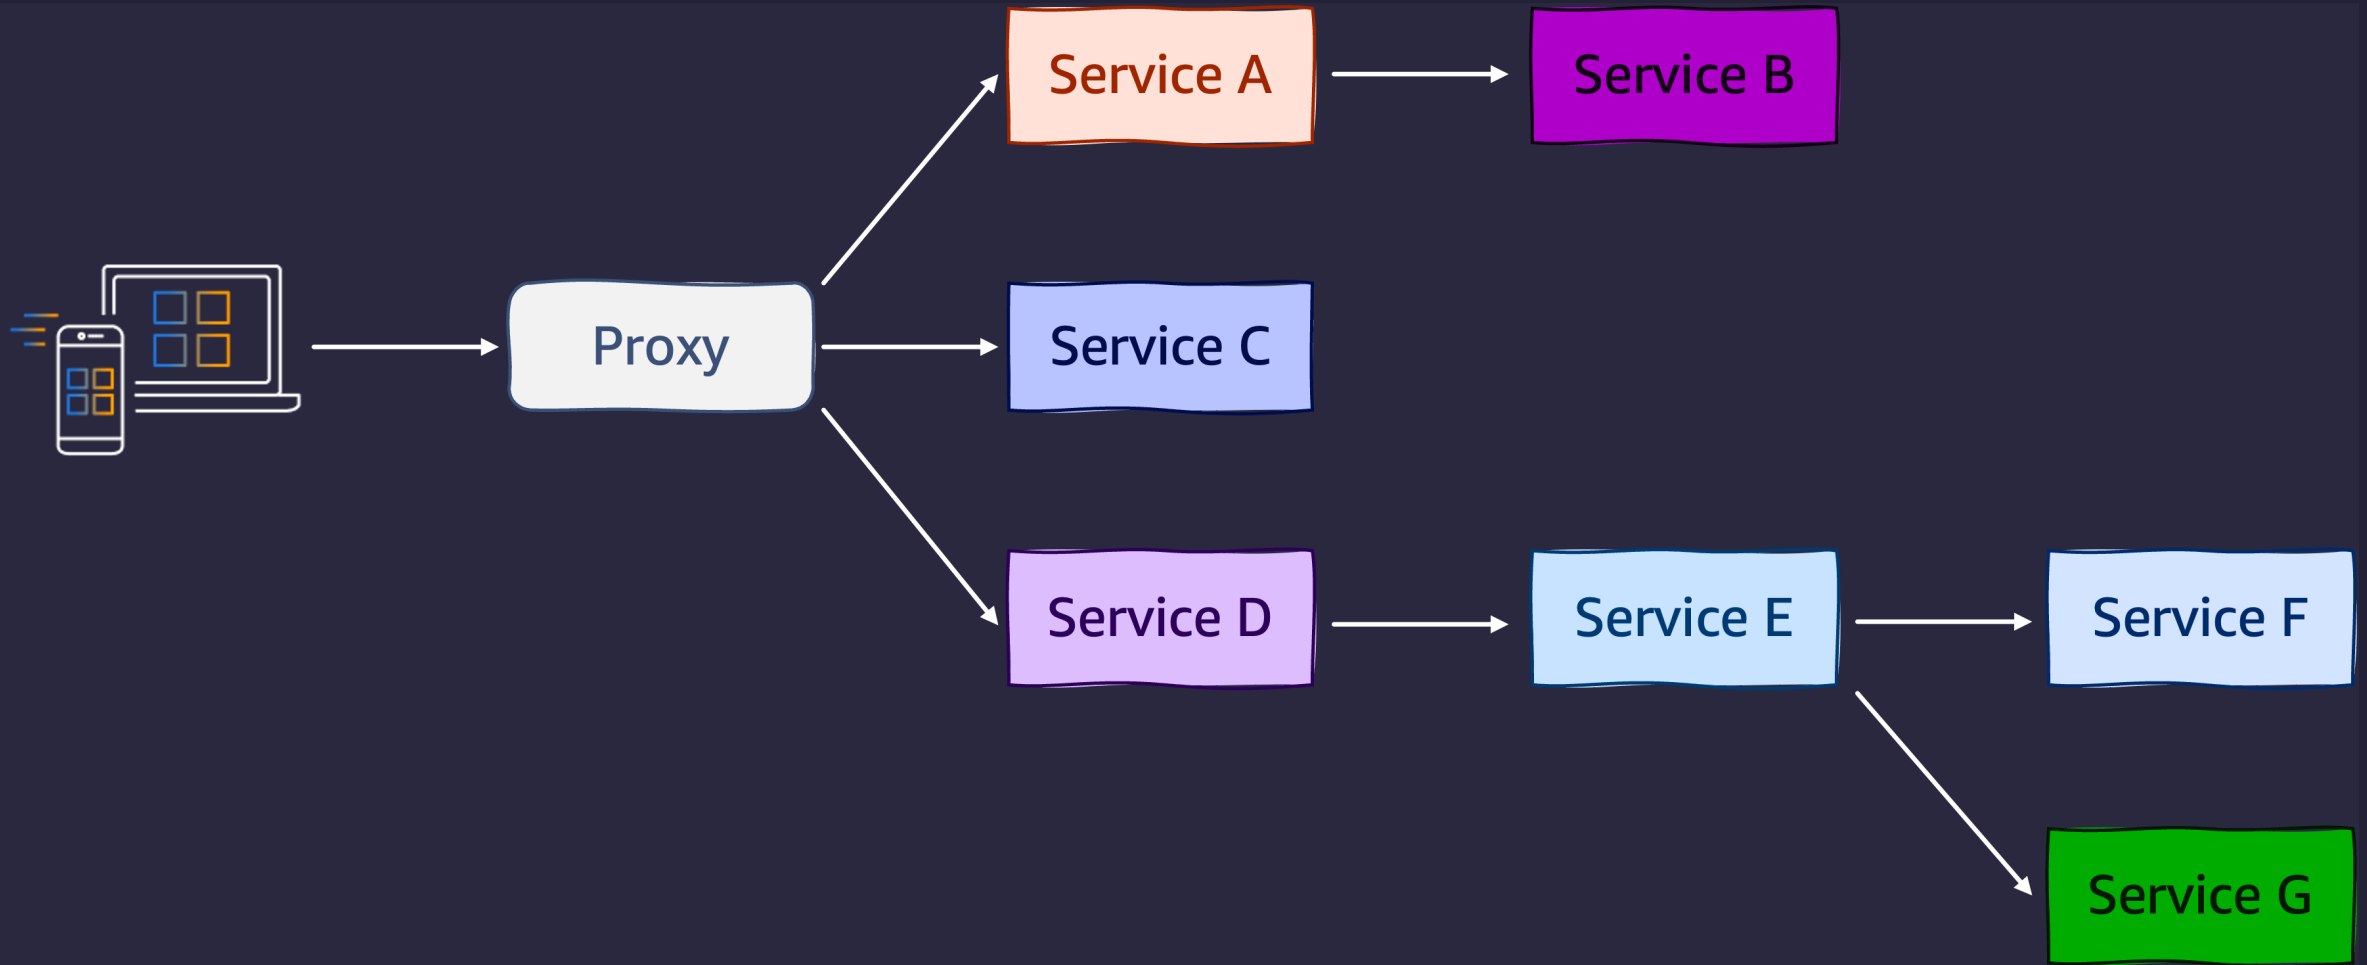
Distributed Trace

A **telemetry** method that indicates **latency** in specific parts of a system's process:

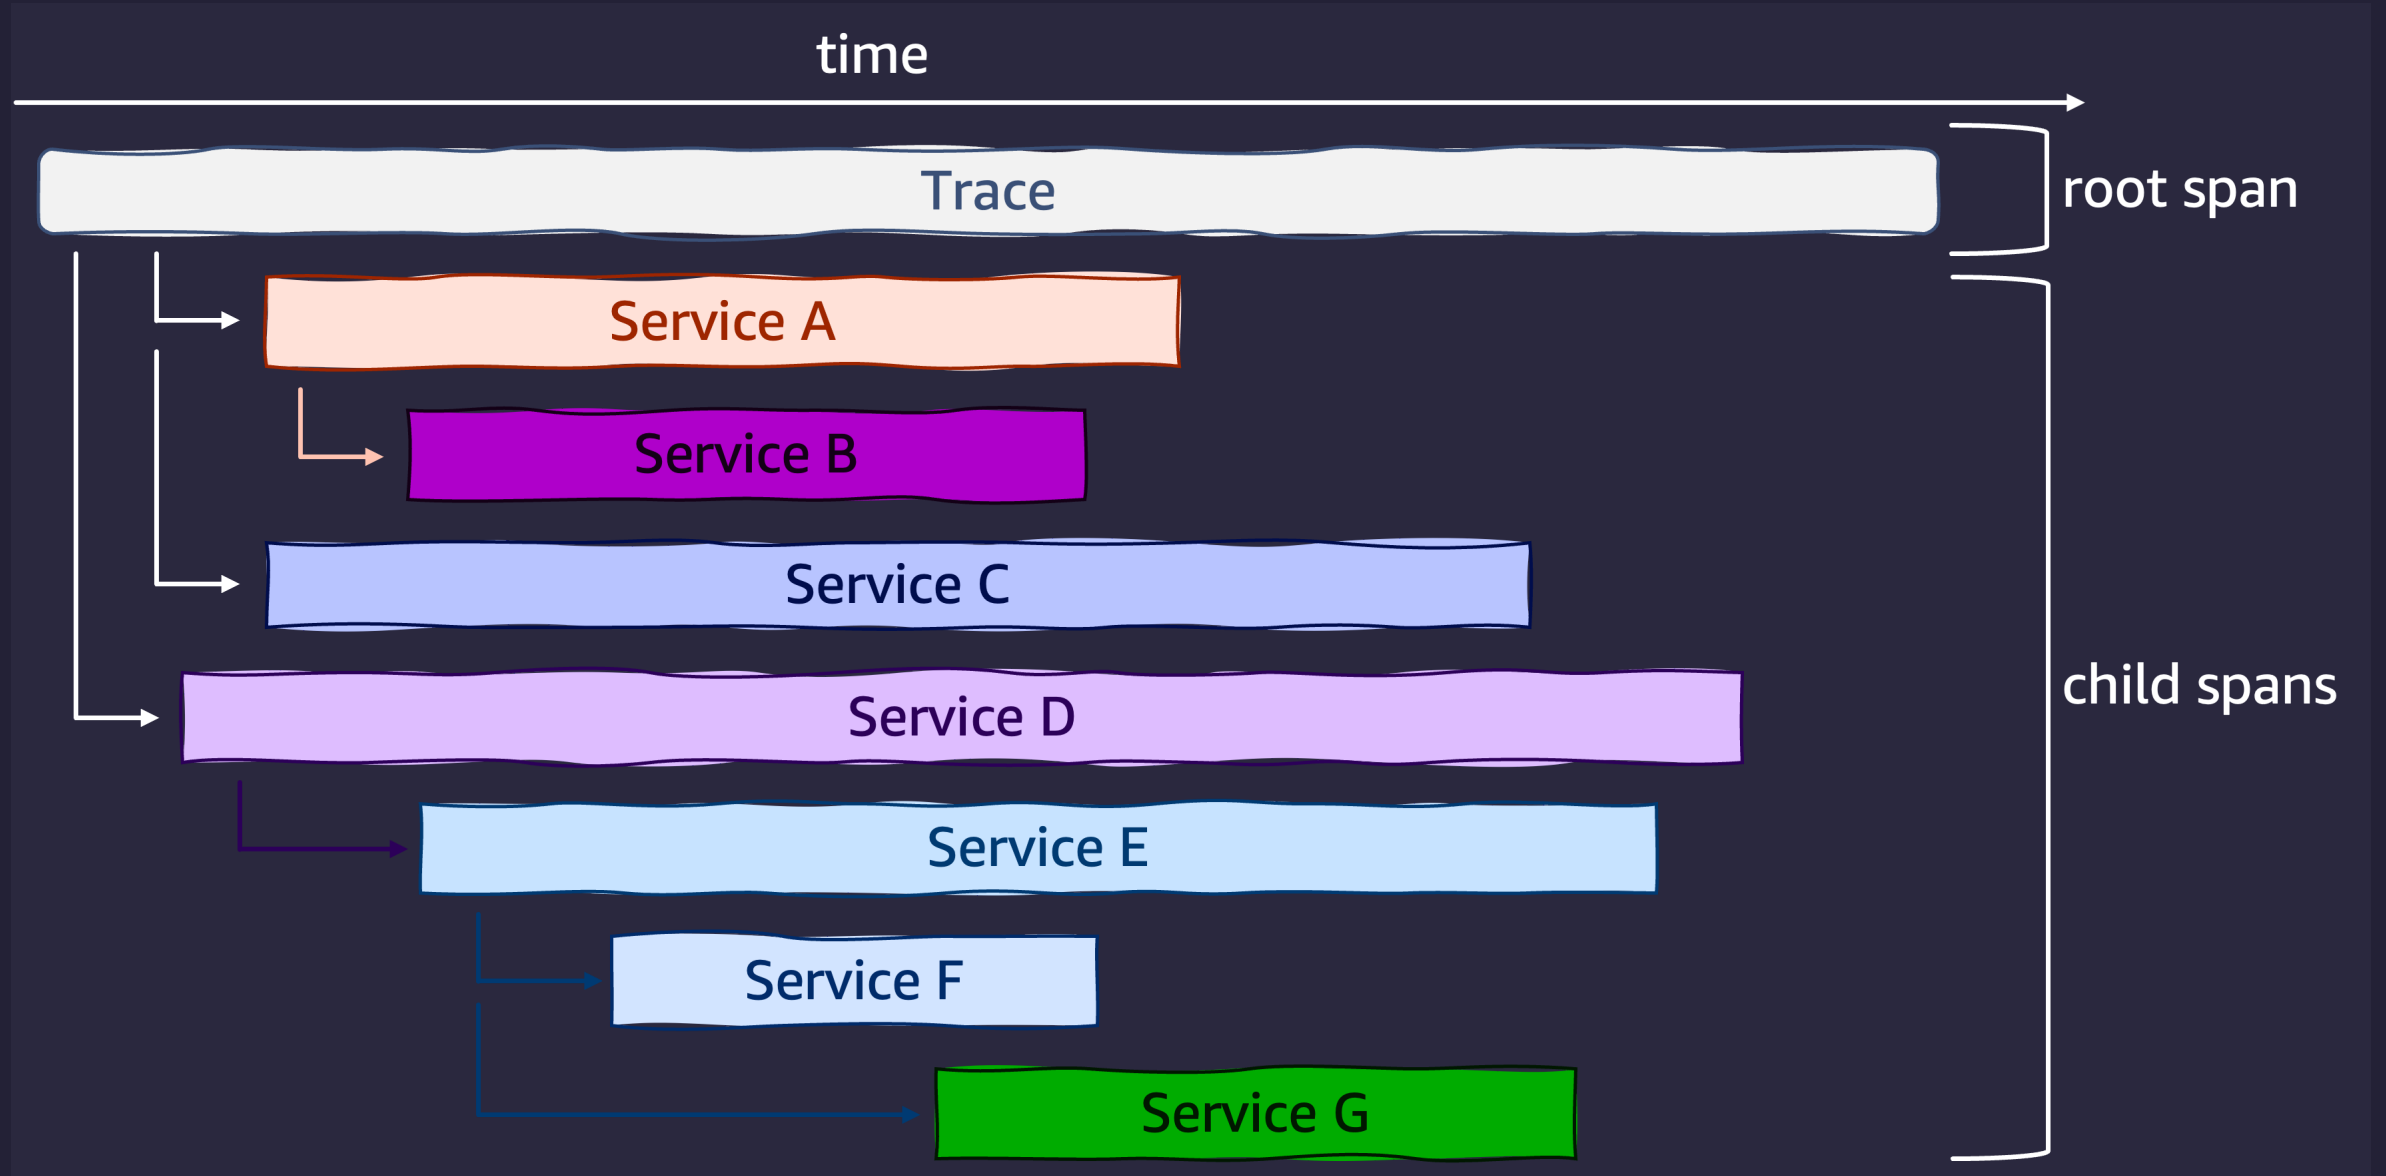
Records the path that requests take as they propagate through microservices and serverless architectures.

It is crucial for **observability** because it measures **dependencies** and **relationships** among numerous components in modern architectures like microservices environments, identifying **latency bottlenecks**.

Trace

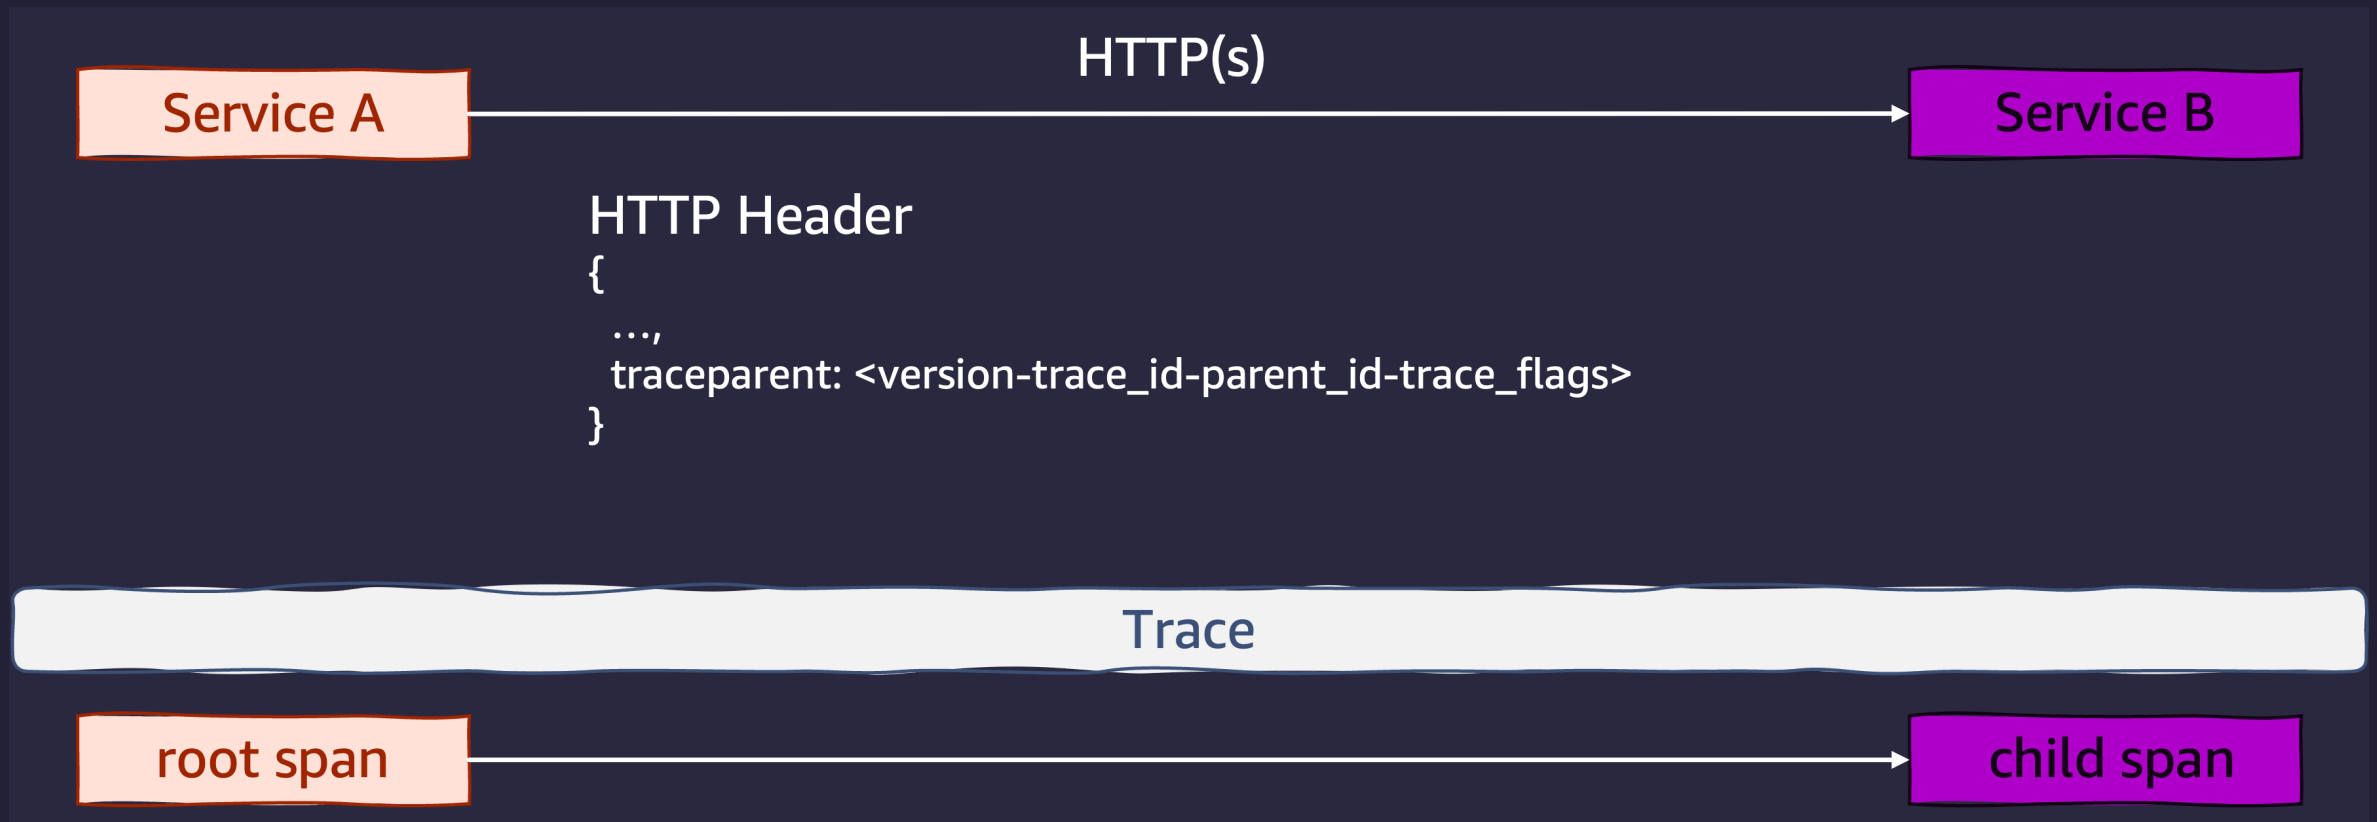


Trace (Span)

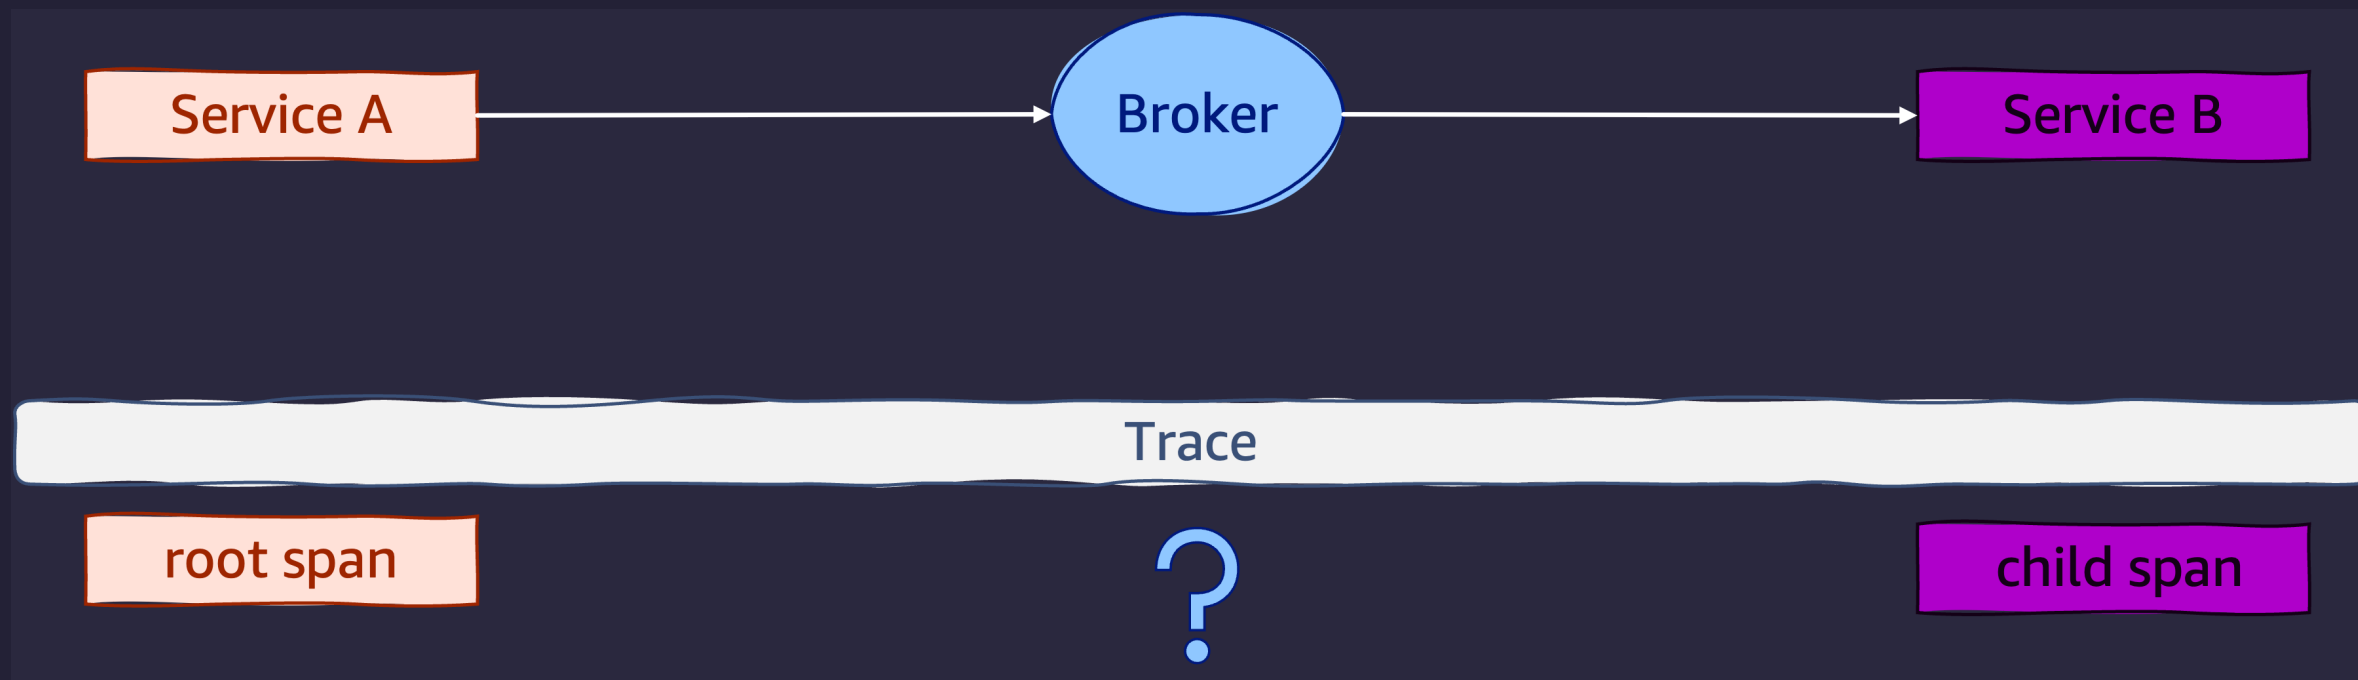


Trace - context propagation

W3C Trace Context : <https://w3c.github.io/trace-context/>



Trace - context propagation



Trace - context propagation



Trace

root span

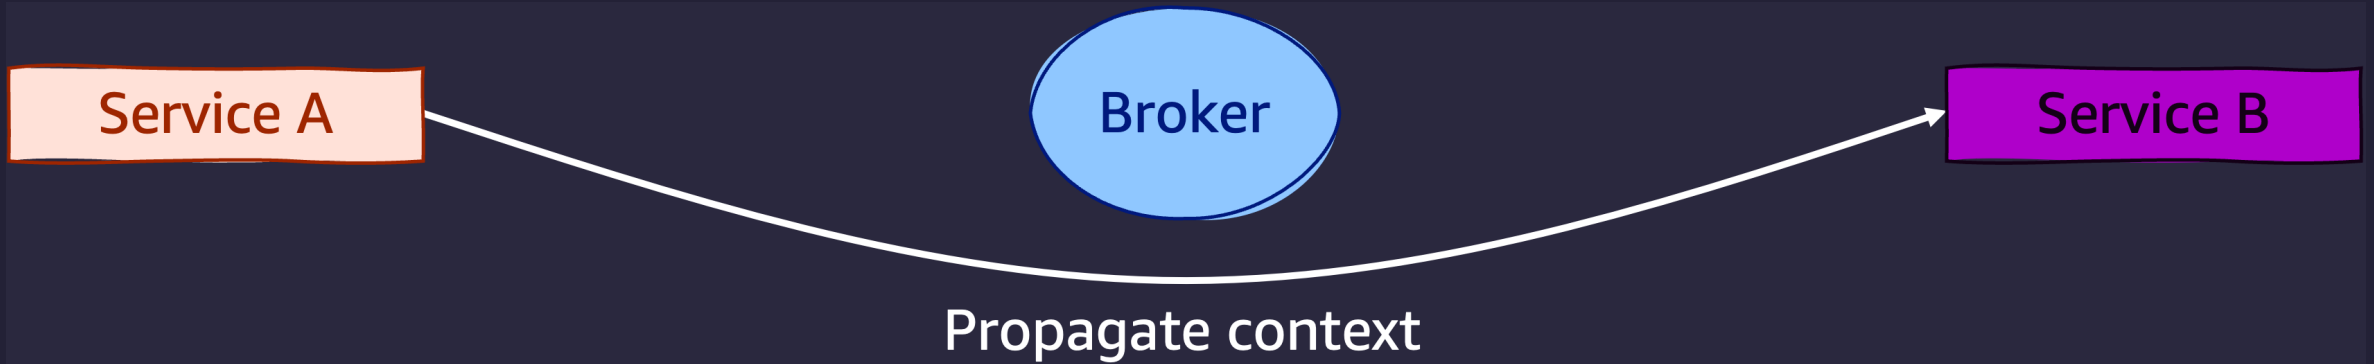
```
{
  "name": "root-span",
  "context": {
    "trace_id": "0x5b8aa5a2d2c872e8321cf37308d69df2",
    "span_id": "0x051581bf3cb55c13"
  }
  "parent_id": null,
  "start_time": "2022-04-29T18:52:58.114201Z",
  "end_time": "2022-04-29T18:52:58.114687Z"
}
```

?

child span

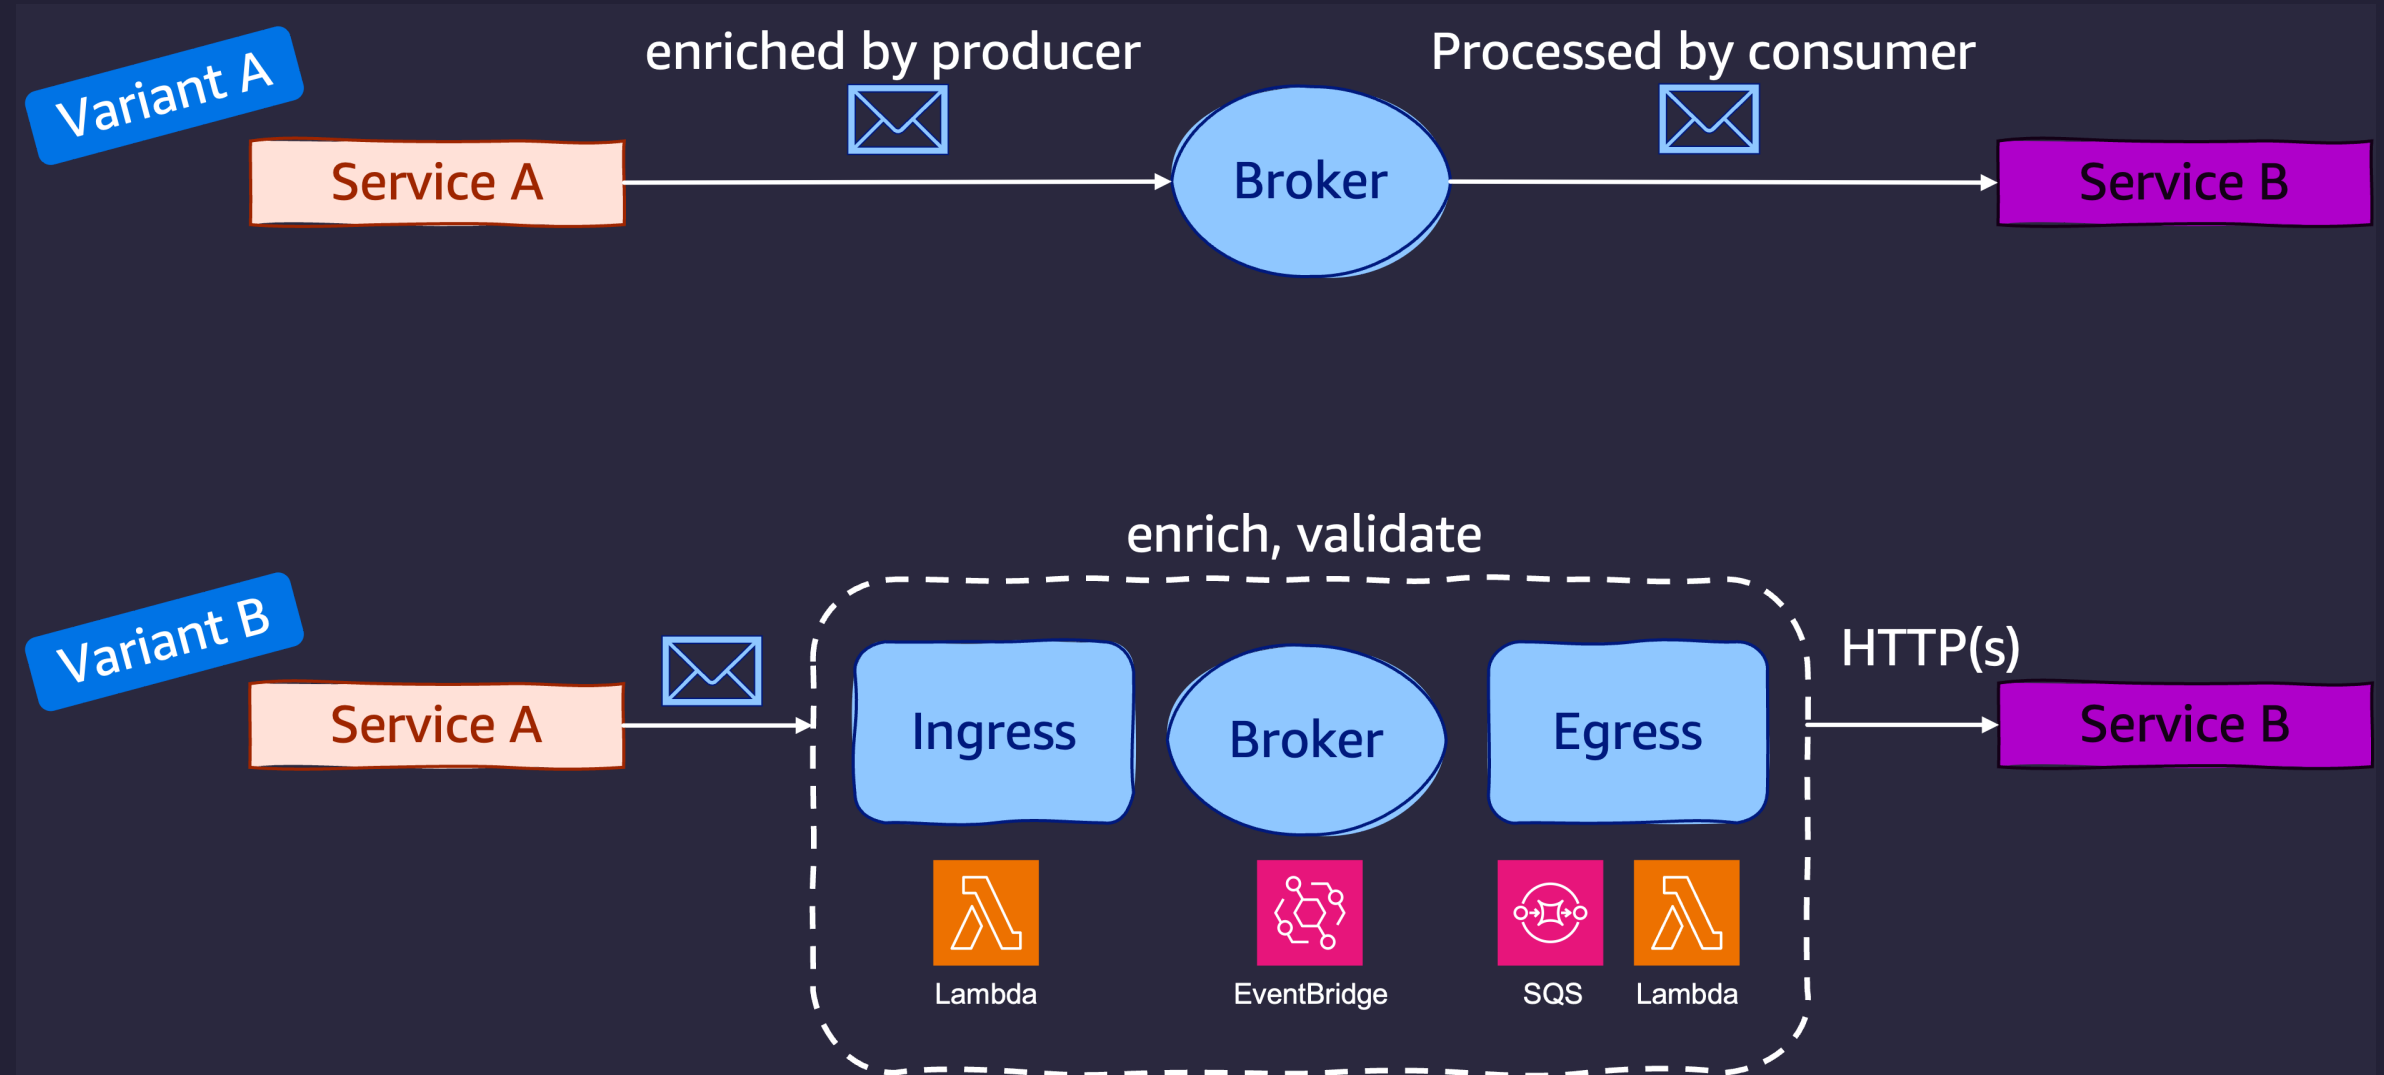
```
{
  "name": "child-span",
  "context": {
    "trace_id": "0x5b8aa5a2d2c872e8321cf37308d69df2",
    "span_id": "0x5fb397be34d26b51"
  }
  "parent_id": "0x051581bf3cb55c13",
  "start_time": "2022-04-29T18:52:58.114504Z",
  "end_time": "2022-04-29T22:52:58.114561Z"
}
```

Trace - context propagation



```
{  
  → "metadata": {  
    < some mandatory fields like timestamps, domain information, ...>  
    "otel": {  
      "trace_parent": {  
        "trace_id": "0x5b8aa...",  
        "span_id": "0x05158..."  
      }  
    }  
  },  
  → "data": {  
    < original producer event >  
  }  
}
```

Trace - context propagation



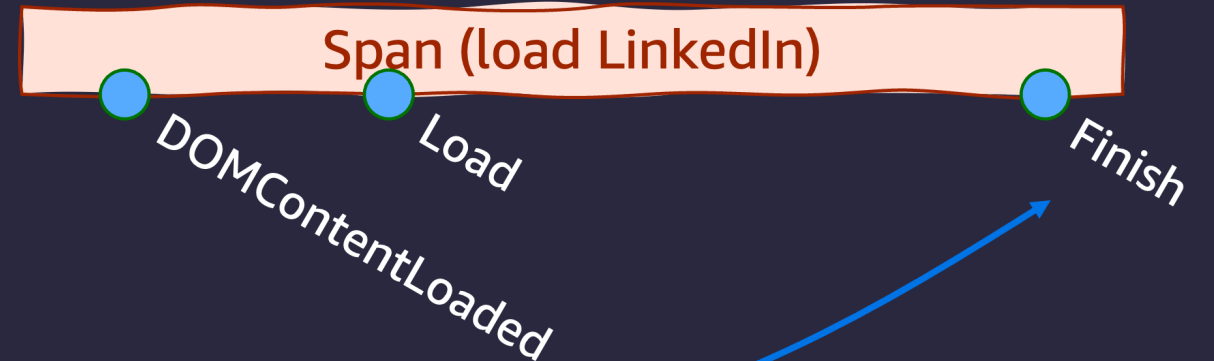
Span

```
{
  "name": "/v1/sys/health",
  "context": {
    "trace_id": "7bba9f33312b3dbb8b2c2c62bb7abe2d",
    "span_id": "086e83747d0e381e"
  },
  "parent_id": "",
  "start_time": "2021-10-22 16:04:01.209458162 +0000 UTC",
  "end_time": "2021-10-22 16:04:01.209514132 +0000 UTC",
  "status_code": "STATUS_CODE_OK",
  "status_message": "",
  "attributes": {
    "http.scheme": "http",
    "http.host": "10.177.2.152:26040",
  },
  "events": [
    {
      "name": "",
      "message": "OK",
      "timestamp": "2021-10-22 16:04:01.209512872 +0000 UTC"
    }
  ]
}
```

Span

170 / 521 requests | 2.5 MB / 2.8 MB transferred | 32.8 MB / 34.9 MB resources | Finish: 18.72 s | DOMContentLoaded: 1.66 s | Load: 2.54 s

```
{
  "attributes": {
    "response_size": "2.5",
    "content_size": "2.8",
  },
  "events": [
    {
      "name": "Finish",
      "message": "OK",
      "timestamp": "2021-10-22 16:04:01.209512872 +0000 UTC"
    }
  ]
}
```

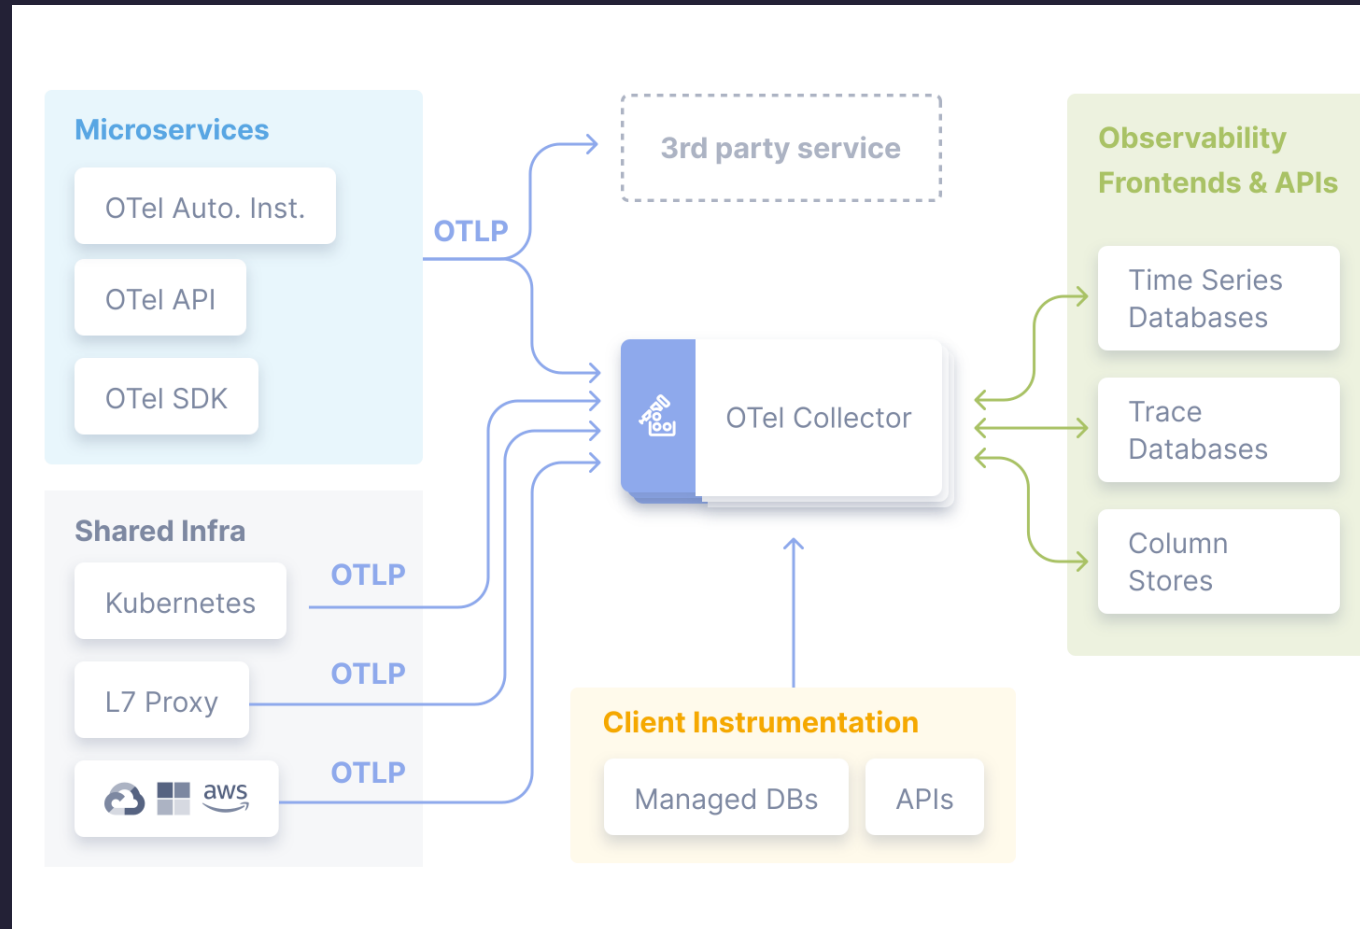


<https://opentelemetry.io/docs/specs/otel/trace/exceptions/>

Hello, OpenTelemetry

Open source project hosted on CNCF Specifications, Implementations for instrumentation and transmissions of telemetry data (metrics, logs, traces)

1. Cross-language specifications
2. OpenTelemetry Collector (agent)
3. SDKs for each language
4. Auto Instrumentation



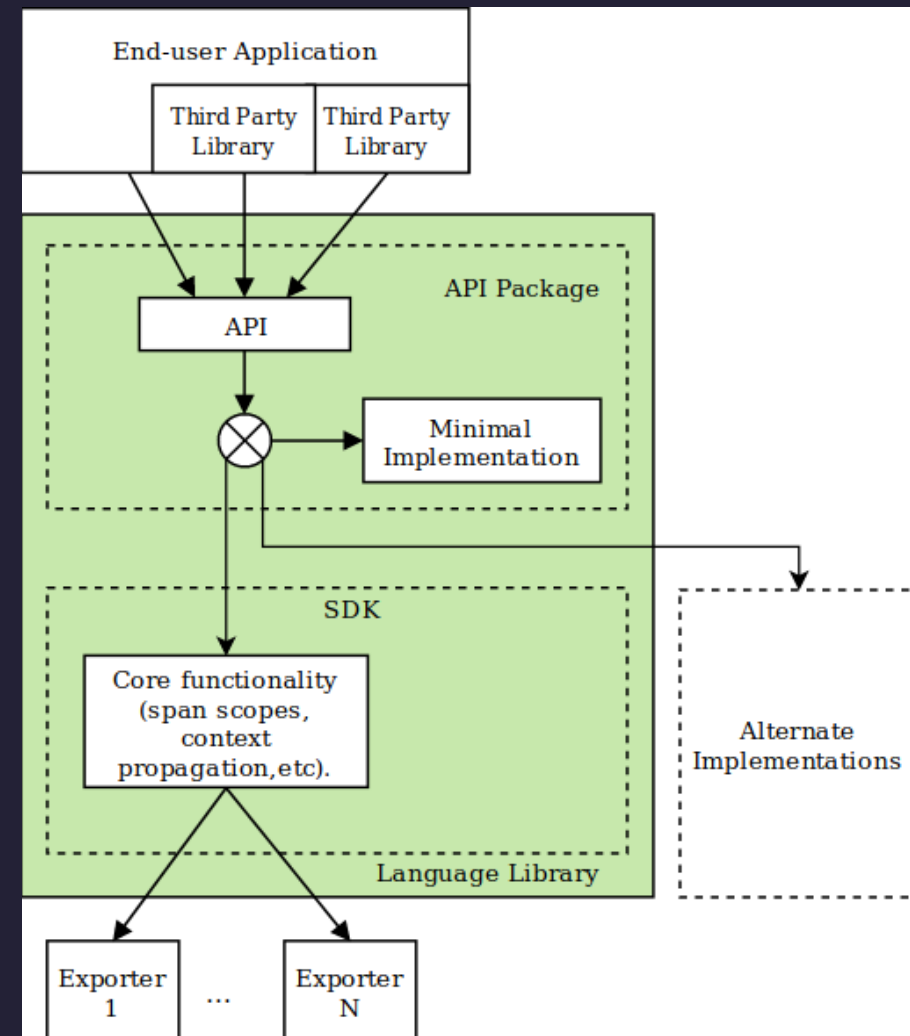
OpenTelemetry Instrumentation

1. Code-based solutions via official APIs and SDKs for most languages

- **API** defines data types and how to generate telemetry data.
- **SDK** defines a language-specific implementation of the API, plus configuration, data processing and exporting.

2. Zero-code solutions

- Go, .NET, PHP, Python, Java, JavaScript



Manually Instrumentation (Python)

```
@app.route("/server_request")
def server_request():
    with tracer.start_as_current_span(
        "server_request",
        context=extract(request.headers),
        kind=trace.SpanKind.SERVER,
        attributes=collect_request_attributes(request.environ),
    ):
        print(request.args.get("param"))
        return "served"
```

Programmatically-instrumented server (Python)

```
instrumentor = FlaskInstrumentor()

app = Flask(__name__)

instrumentor.instrument_app(app)
# instrumentor.instrument_app(app, excluded_urls="/server_request")
@app.route("/server_request")
def server_request():
    print(request.args.get("param"))
    return "served"
```

OpenTelemetry Registry

OpenTelemetry instrumentation libraries

<https://opentelemetry.io/ecosystem/registry/>

Search **NGINX**

The screenshot shows the OpenTelemetry Registry website. The header includes the OpenTelemetry logo and navigation links: Docs, Ecosystem, Status, Community, Blog, English, and a search bar. The main heading is "Registry" with the subtitle "Find libraries, plugins, integrations, and other useful tools for using and extending OpenTelemetry." Below this, a search bar shows "Search 838 entries" for the term "nginx". Two results are displayed:

- NGINX Instrumentation** by OpenTelemetry Authors. Description: "NGINX OpenTelemetry module to add OpenTelemetry distributed tracing support to NGINX." Metadata: C++ Language, Instrumentation Component, Apache 2.0 License. Link: Repository.
- NGINX Native OTel Module** by NGINX, Inc. Metadata: C++ Language, Instrumentation Component, Apache-2.0 License. Badges: new, first party integration. Links: Website, Repository.

OpenTelemetry Protocol (OTLP)

<https://github.com/open-telemetry/opentelemetry-proto/blob/main/docs/specification.md>

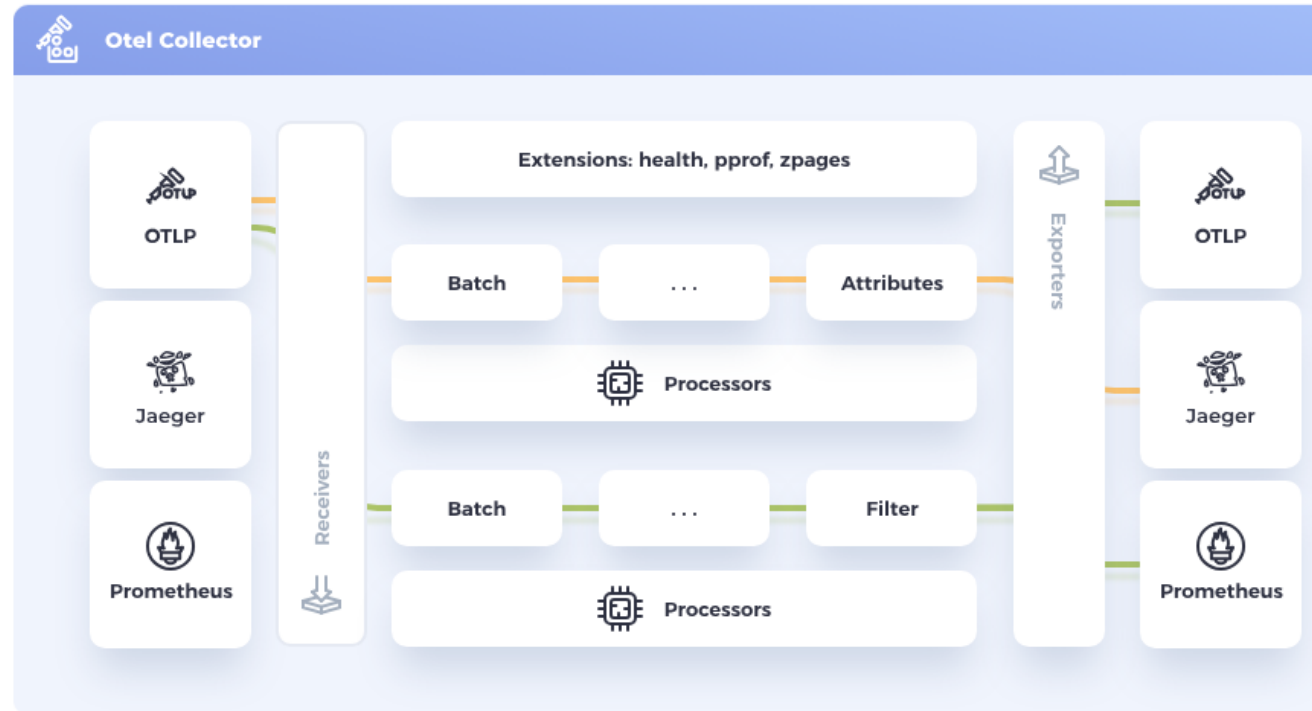
OTLP is implemented over **gRPC** and **HTTP** transports and specifies the Protocol Buffers schema used for payloads.

OTLP is a request/response style protocol where **clients** send requests and the **server** replies with corresponding responses.

All server components must support the following transport compression options: **none**, **gzip**

OpenTelemetry Collector

<https://opentelemetry.io/docs/collector/configuration>



OTEL COLLECTOR

OpenTelemetry Collector Receiver

```
receivers:
  jaeger:
    protocols:
      grpc:
        endpoint: 0.0.0.0:14250
      thrift_compact:
      thrift_http:
  kafka:
    protocol_version: 2.0.0
  prometheus:
    config:
      scrape_configs:
        - job_name: otel-collector
          scrape_interval: 5s
          static_configs:
            - targets: [localhost:8888]
  otlp:
    protocols:
      grpc:
        endpoint: 0.0.0.0:4317
      http:
        endpoint: 0.0.0.0:4318
```

OpenTelemetry Collector Processors

```
processors:  
  attributes:  
    actions:  
      - key: environment  
        value: production  
        action: insert  
      - key: db.statement  
        action: delete  
      - key: email  
        action: hash  
  probabilistic_sampler:  
    hash_seed: 22  
    sampling_percentage: 15  
  memory_limiter:  
    check_interval: 5s  
    limit_mib: 4000  
    spike_limit_mib: 500  
  filter:  
    metrics:  
      include:  
        match_type: regexp  
        metric_names:  
          - prefix/.  
          - prefix_.
```

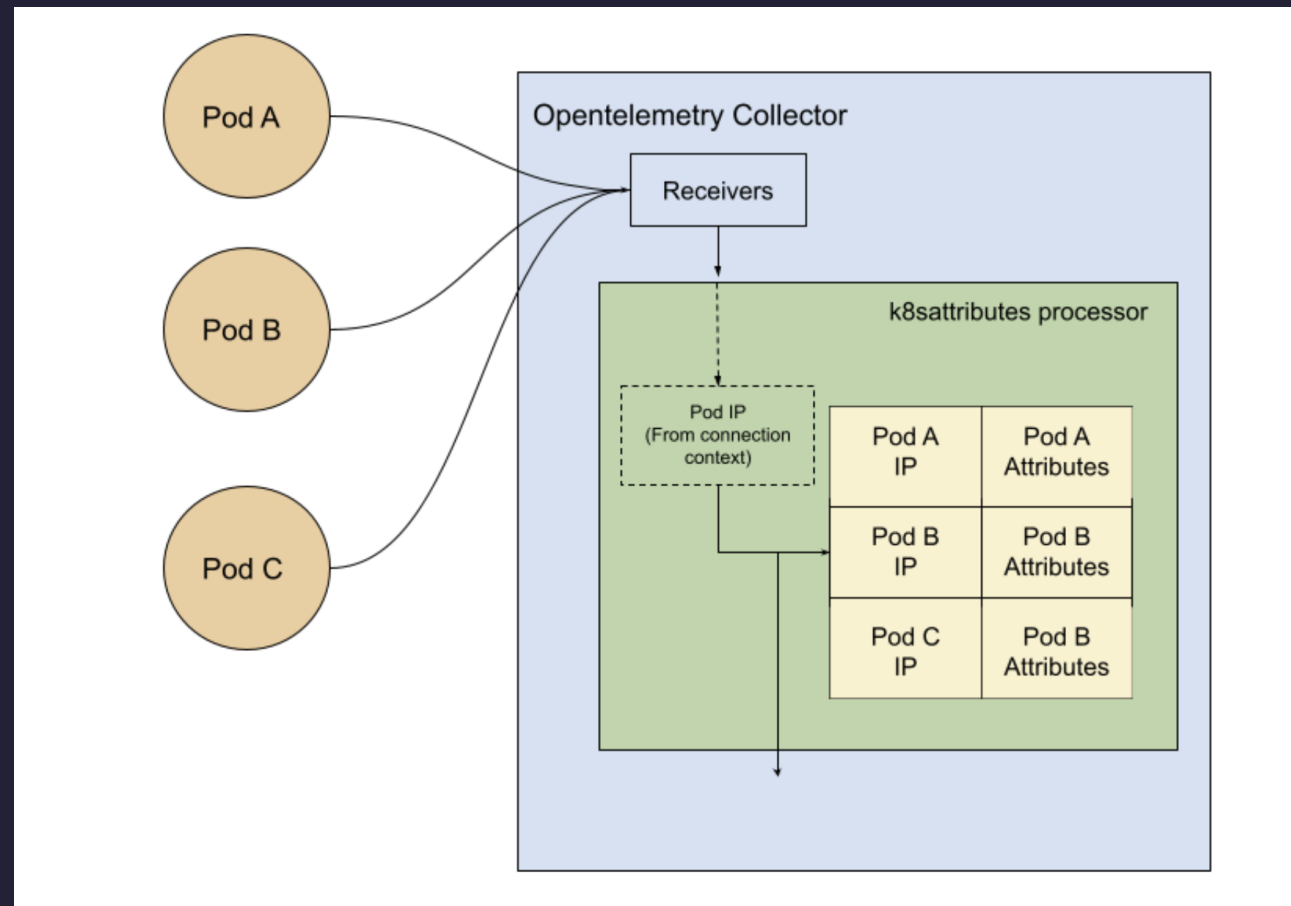
OpenTelemetry Collector Exporter

```
exporters:  
  file:  
    path: ./filename.json  
  jaeger:  
    endpoint: http://jaeger-all-in-one:14250  
    insecure: true  
  kafka:  
    protocol_version: 2.0.0  
  otlphttp:  
    endpoint: https://otlp.example.com:4318  
  prometheus:  
    endpoint: prometheus:8889  
    namespace: default  
  prometheusremotewrite:  
    endpoint: "http://some.url:9411/api/prom/push"  
  zipkin:  
    endpoint: "http://localhost:9411/api/v2/spans"
```


Resource Semantic Conventions

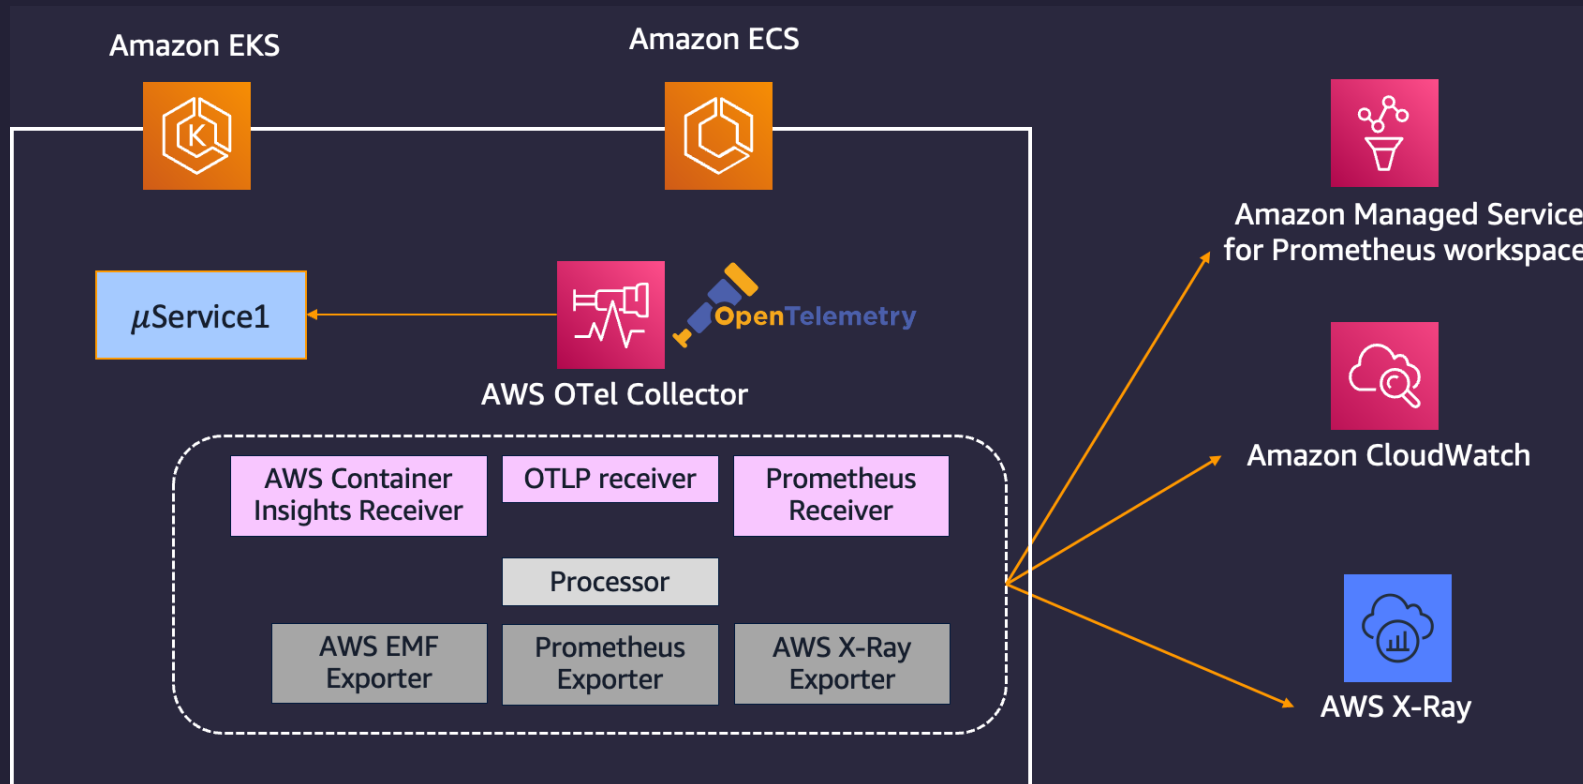
A **Resource** represents the entity producing telemetry as resource attributes.

<https://opentelemetry.io/docs/specs/semconv/>

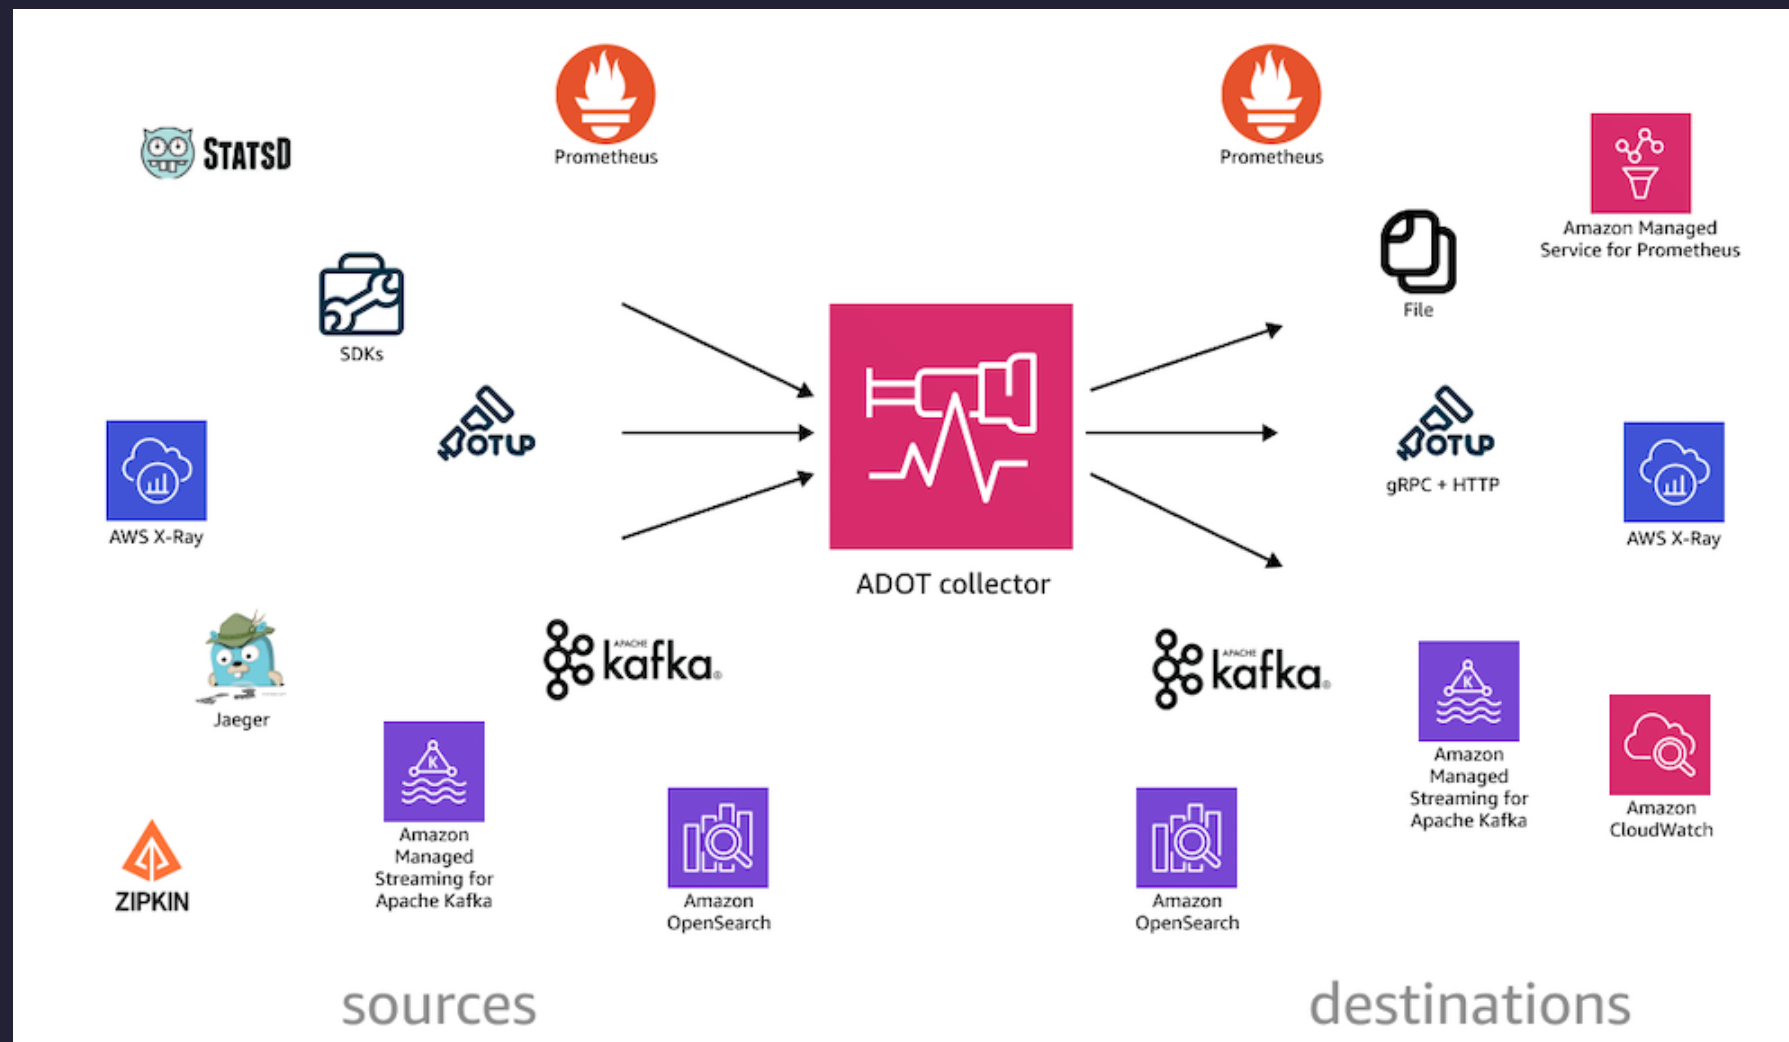


AWS Distro for OpenTelemetry (ADOT)

- Secure, production ready, and supported by AWS OpenTelemetry distribution
- Extend functionality for ease of use on AWS



AWS Distro for OpenTelemetry (ADOT)



OpenTelemetry Ecosystem

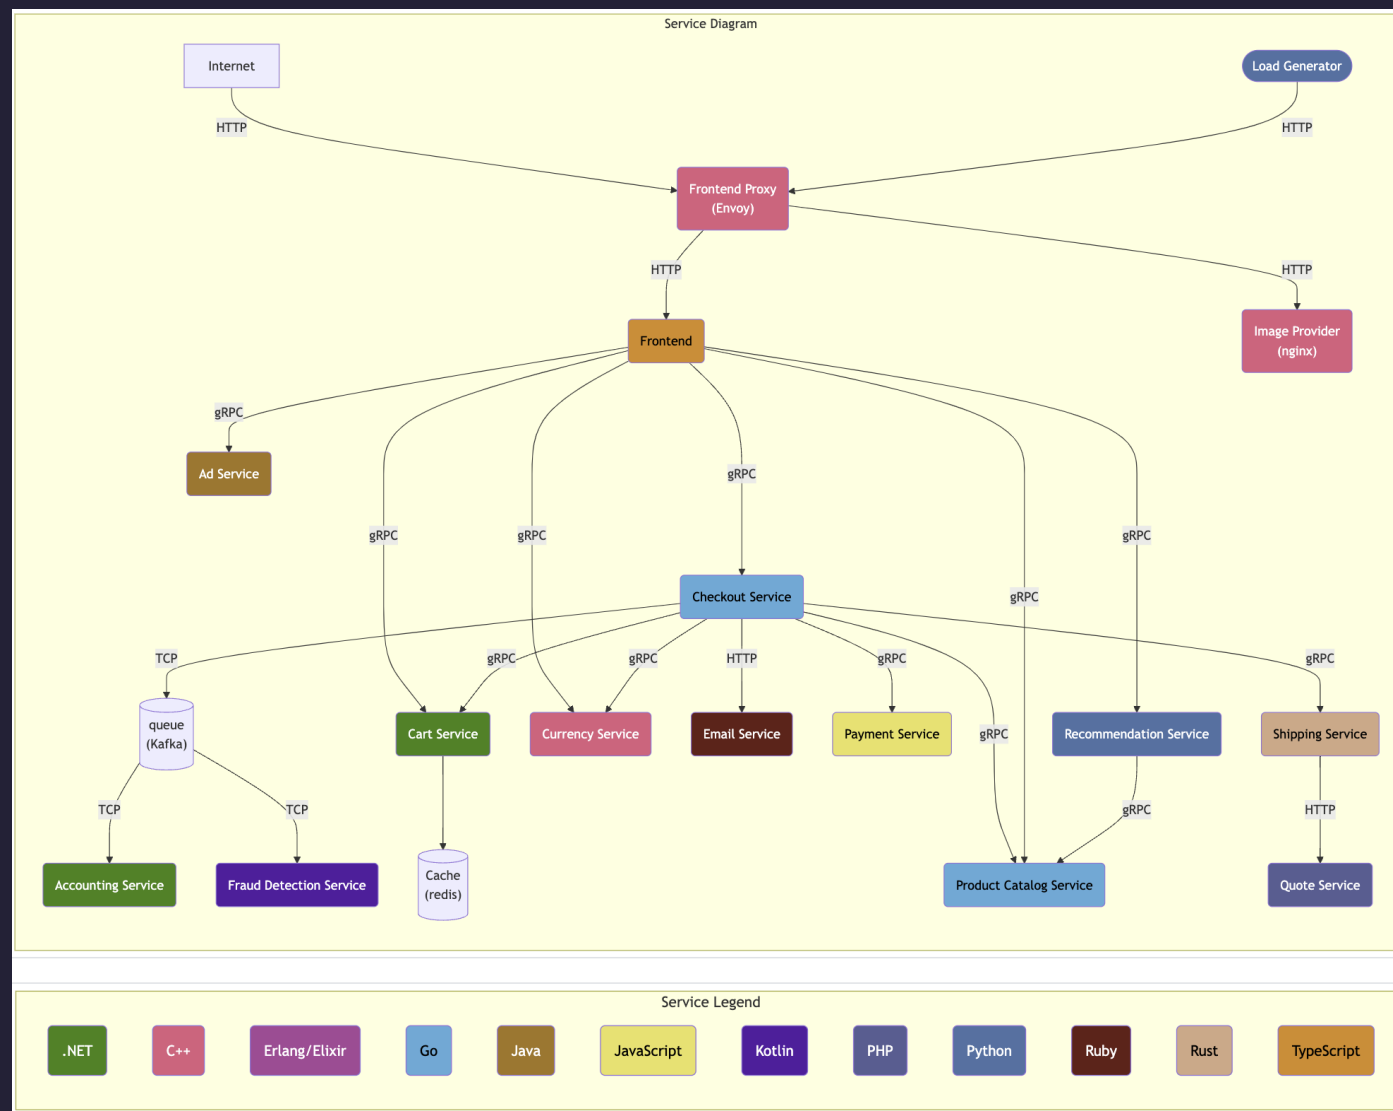
<https://opentelemetry.io/ecosystem/>

- OpenTelemetry Demo
- Registry
- Adopters
Organizations that use OpenTelemetry
- Distributions
List of open source OpenTelemetry distributions maintained by third parties.
- Integrations
Libraries, services, and apps with first-party support for OpenTelemetry.
- Vendors
Vendors who natively support OpenTelemetry

OpenTelemetry Demo

OpenTelemetry Demo is composed of microservices written in different programming languages that talk to each other over gRPC and HTTP; and a load generator which uses Locust to fake user traffic.

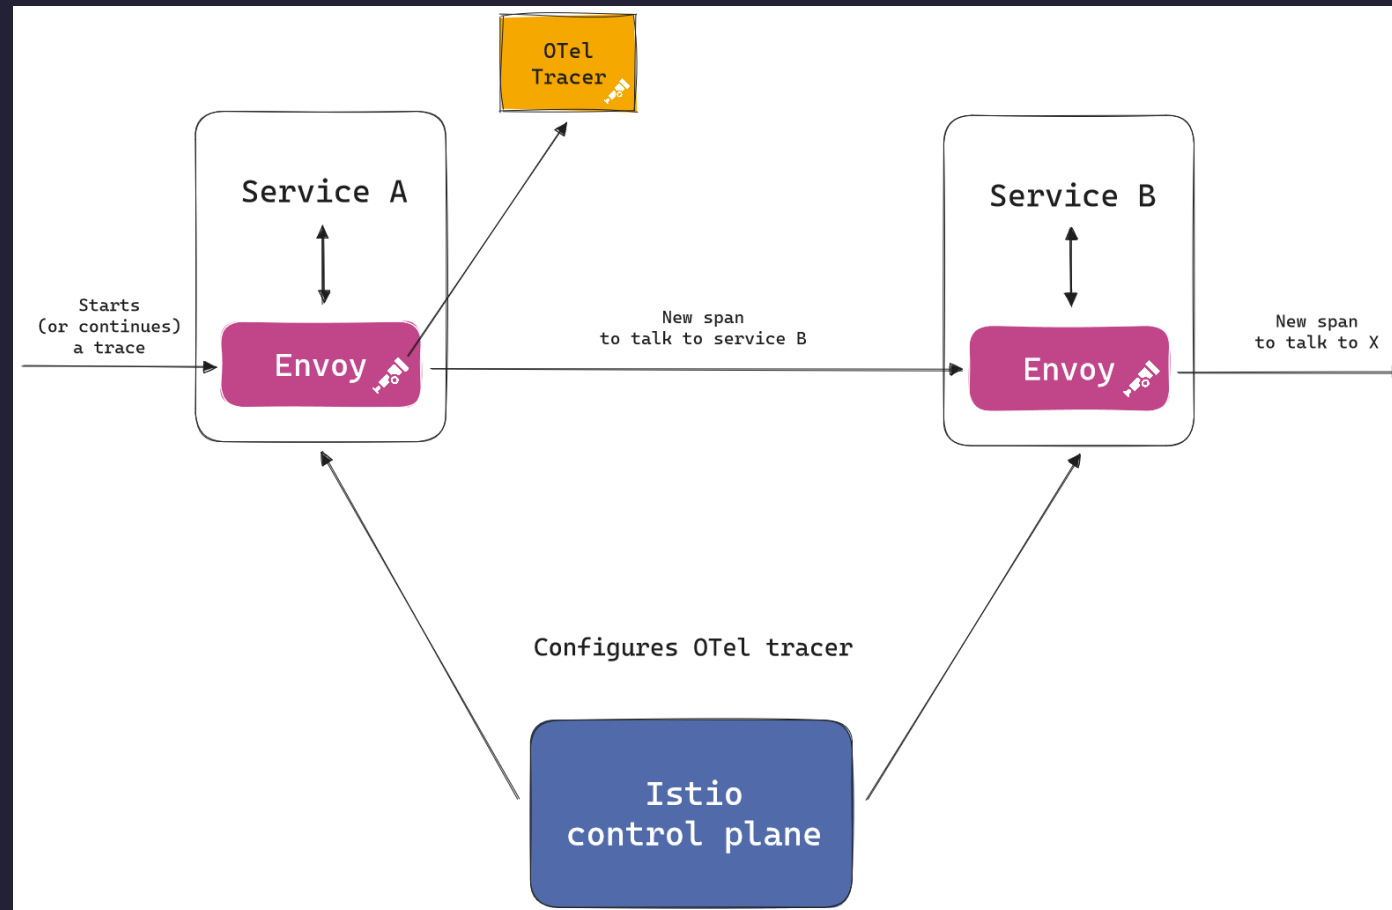
- Web store
- Grafana
- Load Generator
- Jaeger UI



New Otel Feature

Observability in Envoy and Istio

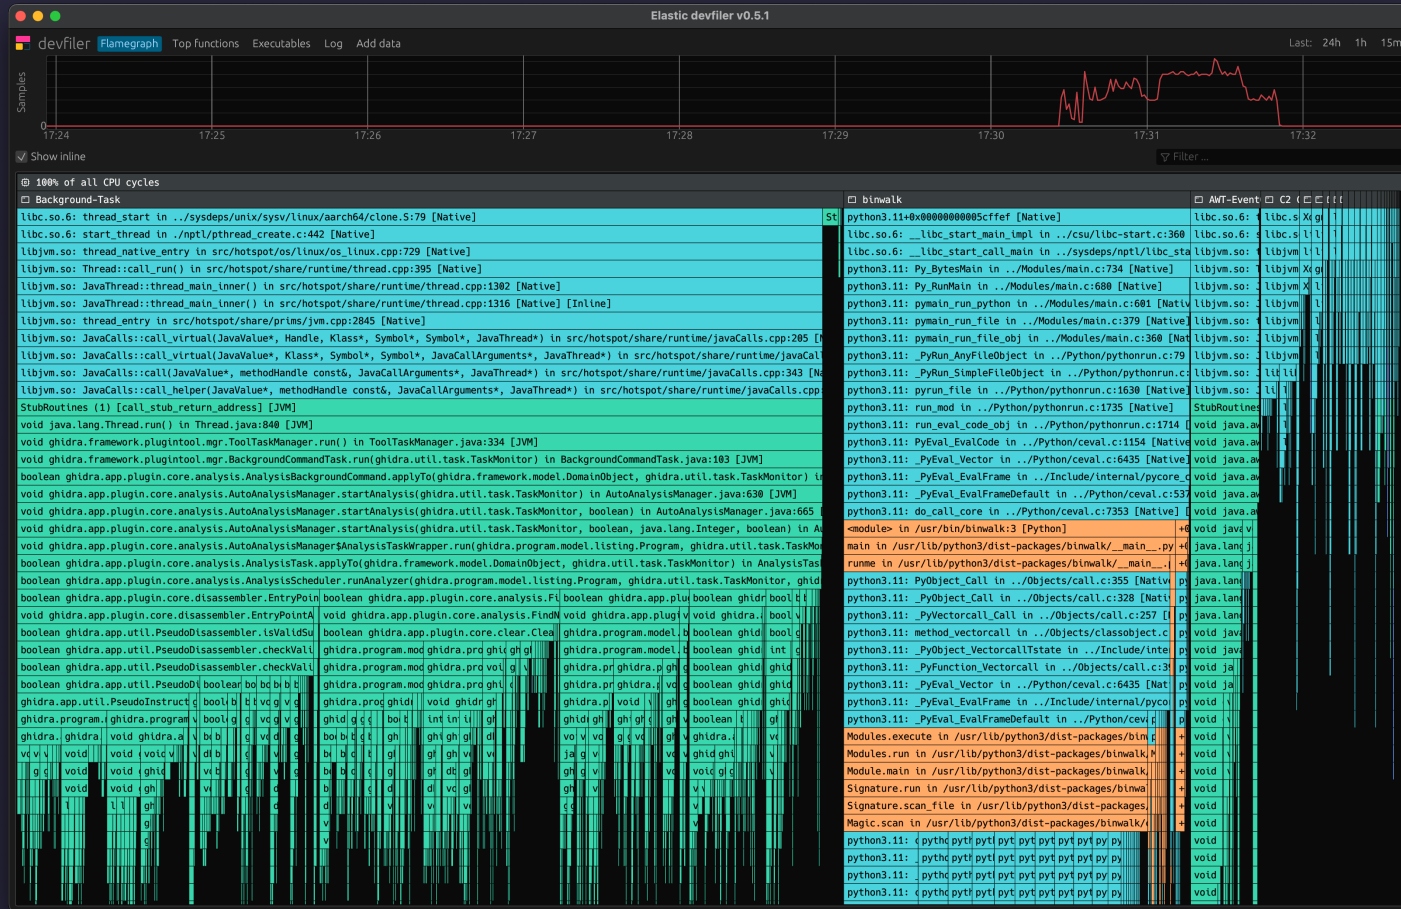
<https://opentelemetry.io/blog/2024/new-otel-features-envoy-istio/>



New Otel Feature

Continuous Profiling Agent

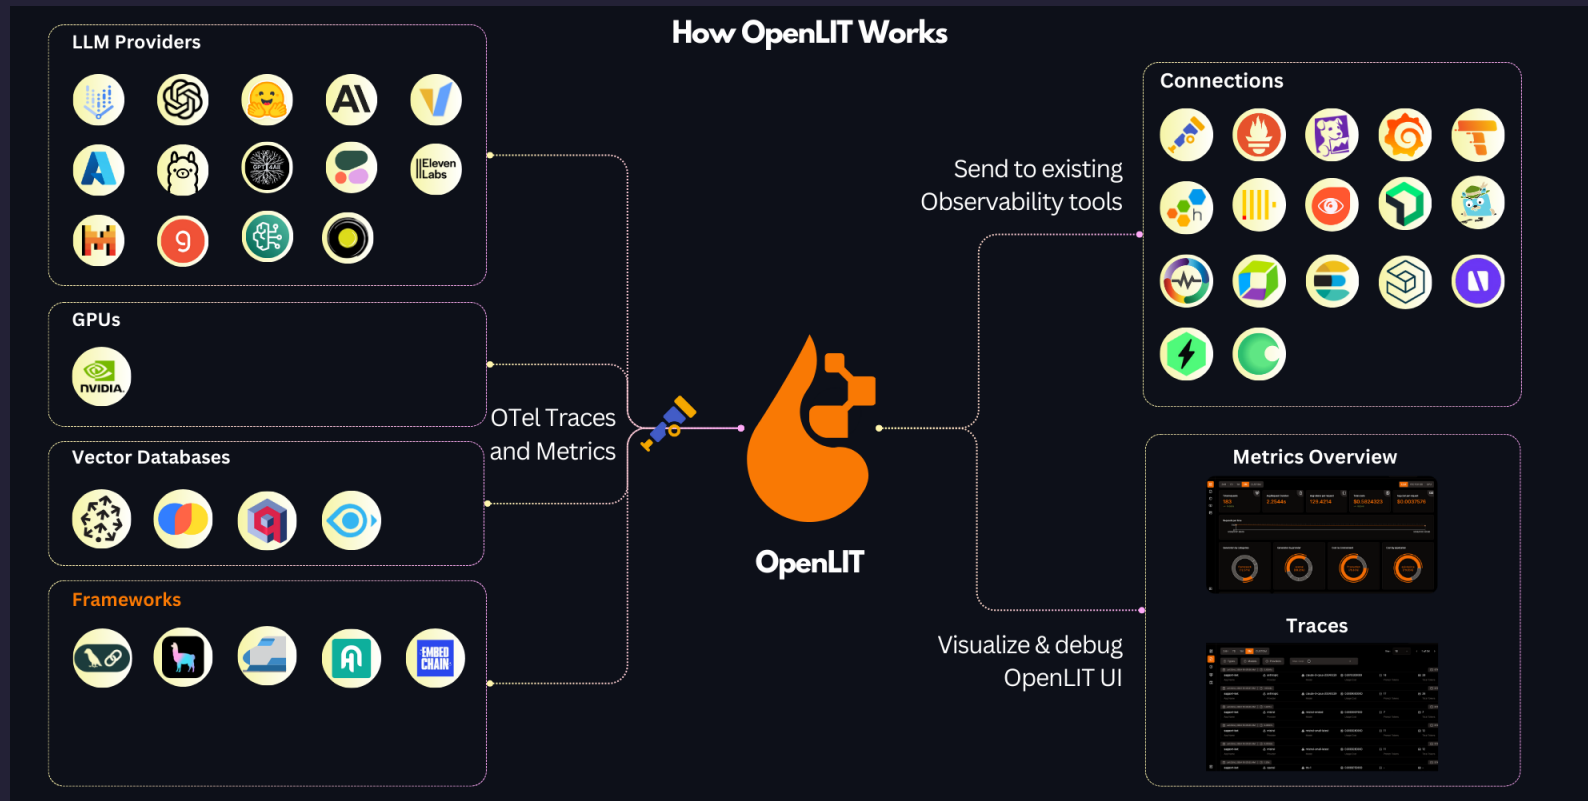
<https://github.com/open-telemetry/opentelemetry-ebpf-profiler>



New Otel Feature

LLM Observability

<https://github.com/openlit/openlit>



References

- <https://opentelemetry.io/docs/>
- <https://w3c.github.io/trace-context/>
- <https://github.com/open-telemetry/opentelemetry-specification>
- <https://opentelemetry.io/docs/specs/semconv/>
- <https://opentelemetry.io/docs/specs/otel/protocol/>
- <https://opentelemetry.io/docs/concepts/sampling/>
- <https://opentelemetry.io/docs/demo/>
- <https://opentelemetry.io/blog/2024/>

Thank You

@ddiwoong

jinwoong@amazon.com