

OpenTelemetry Journey

2024.09

Jinwoong Kim

Who am I?

- 김진웅 (Jinwoong Kim)
- Cloud Architect @AWS Professional Services
- Speaker, Translator
- @ddiwoong

Observability Introduction

A *system* is **observable** if you can determine the *behavior* of the system based on its *outputs*.

Observability

Logs - 무슨 일이 일어난거야? (Lines of text)

```
hikari-pool-1 - Connection is not available, request timed out after 30000ms
```

Metrics - 어떤 지표가 문제야? (Time-ordered set of data)

```
hikaricp_connections_timeout_total{pool="HikariPool-1",} 10.0
```

Tracing - 그 일이 왜 일어났지? (Correlation analysis with Context)

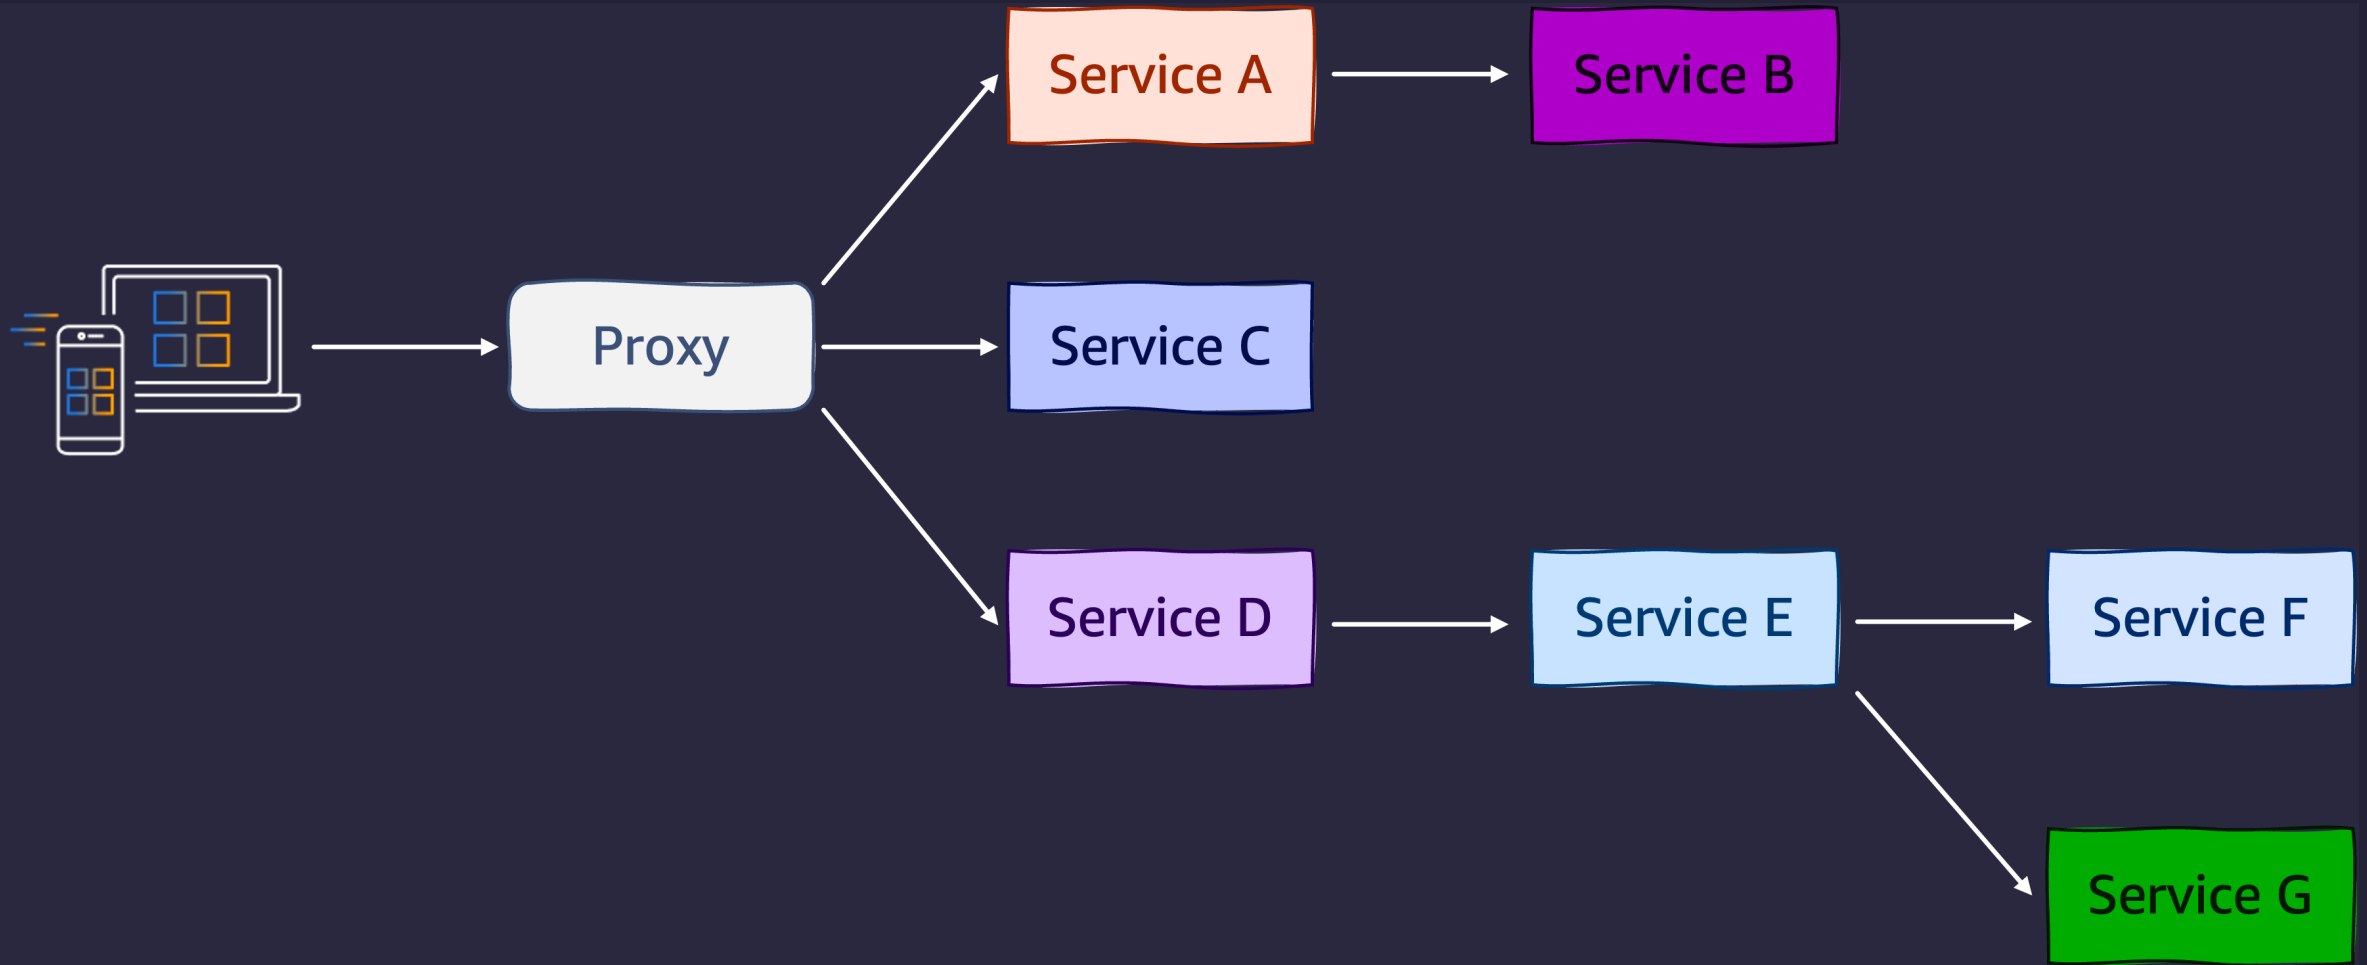
```
2022-05-28 18:09:04.165 INFO [service-b,757d0493f099b94b,4e8d66a6aa1c1ed6] 9989 ---
```

```
[nio-8686-exec-3] c.example.msaerrorresponse.BServiceApi : =====b-service=====
```

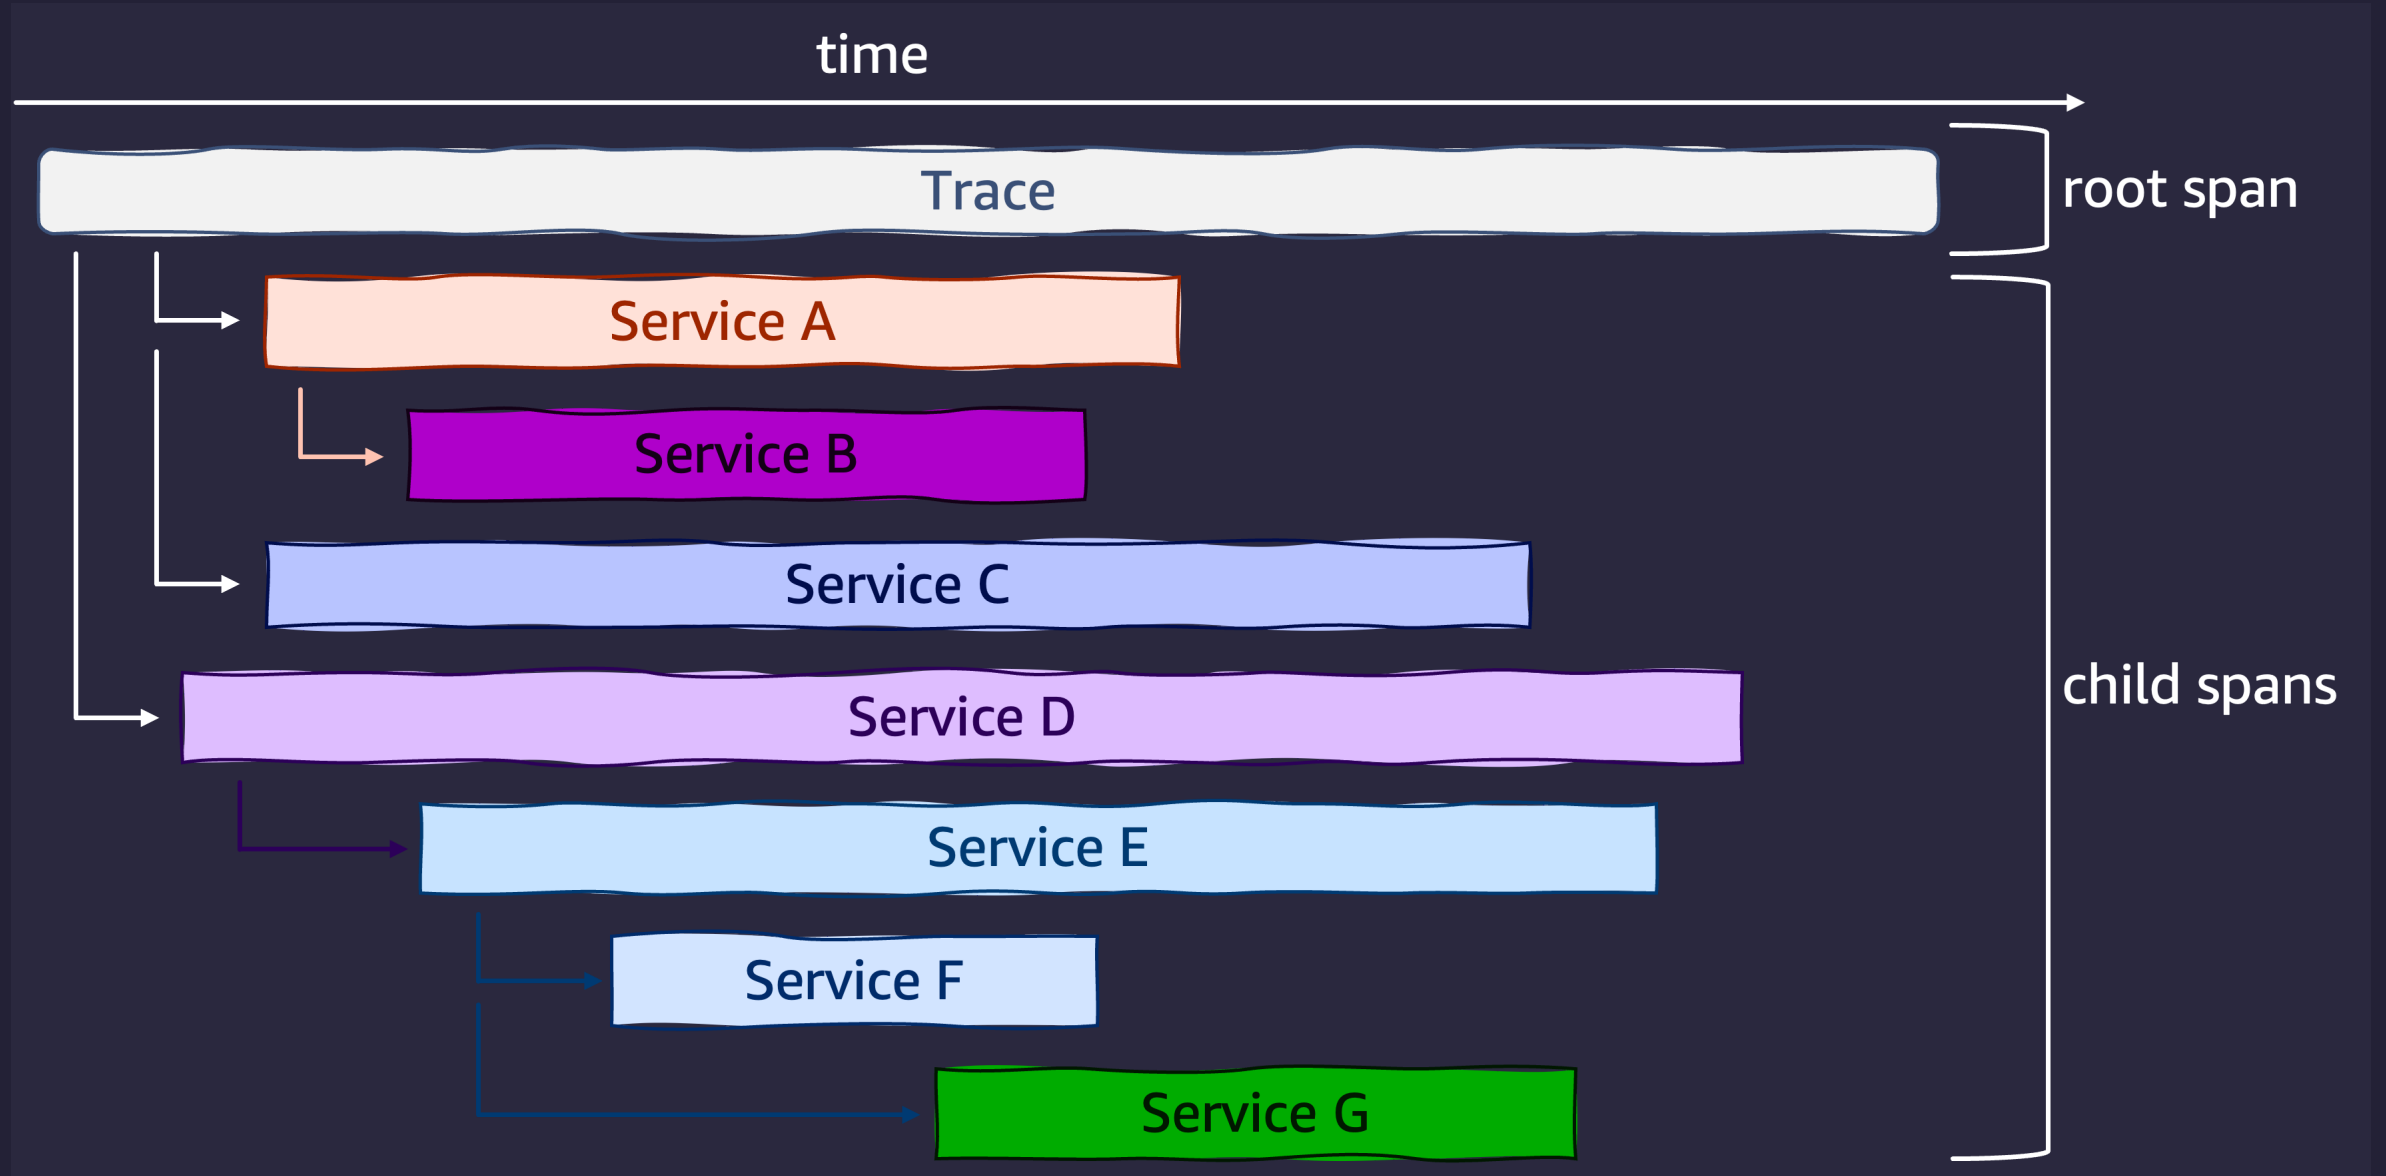
Distributed Trace

- 시스템의 프로세스에서 특정 부분의 지연(latency)을 알려주는 원격 측정(telemetry)방법
- 요청(Requests)이 마이크로서비스 및 서버리스 아키텍처를 통해 전파될때 이동하는 경로를 기록
- 마이크로서비스 환경같은 최신 아키텍처에서 수많은 구성 요소간의 종속성과 관계를 측정하고 지연 병목을 찾아내는 도구이기 때문에 Observability에서 매우 중요

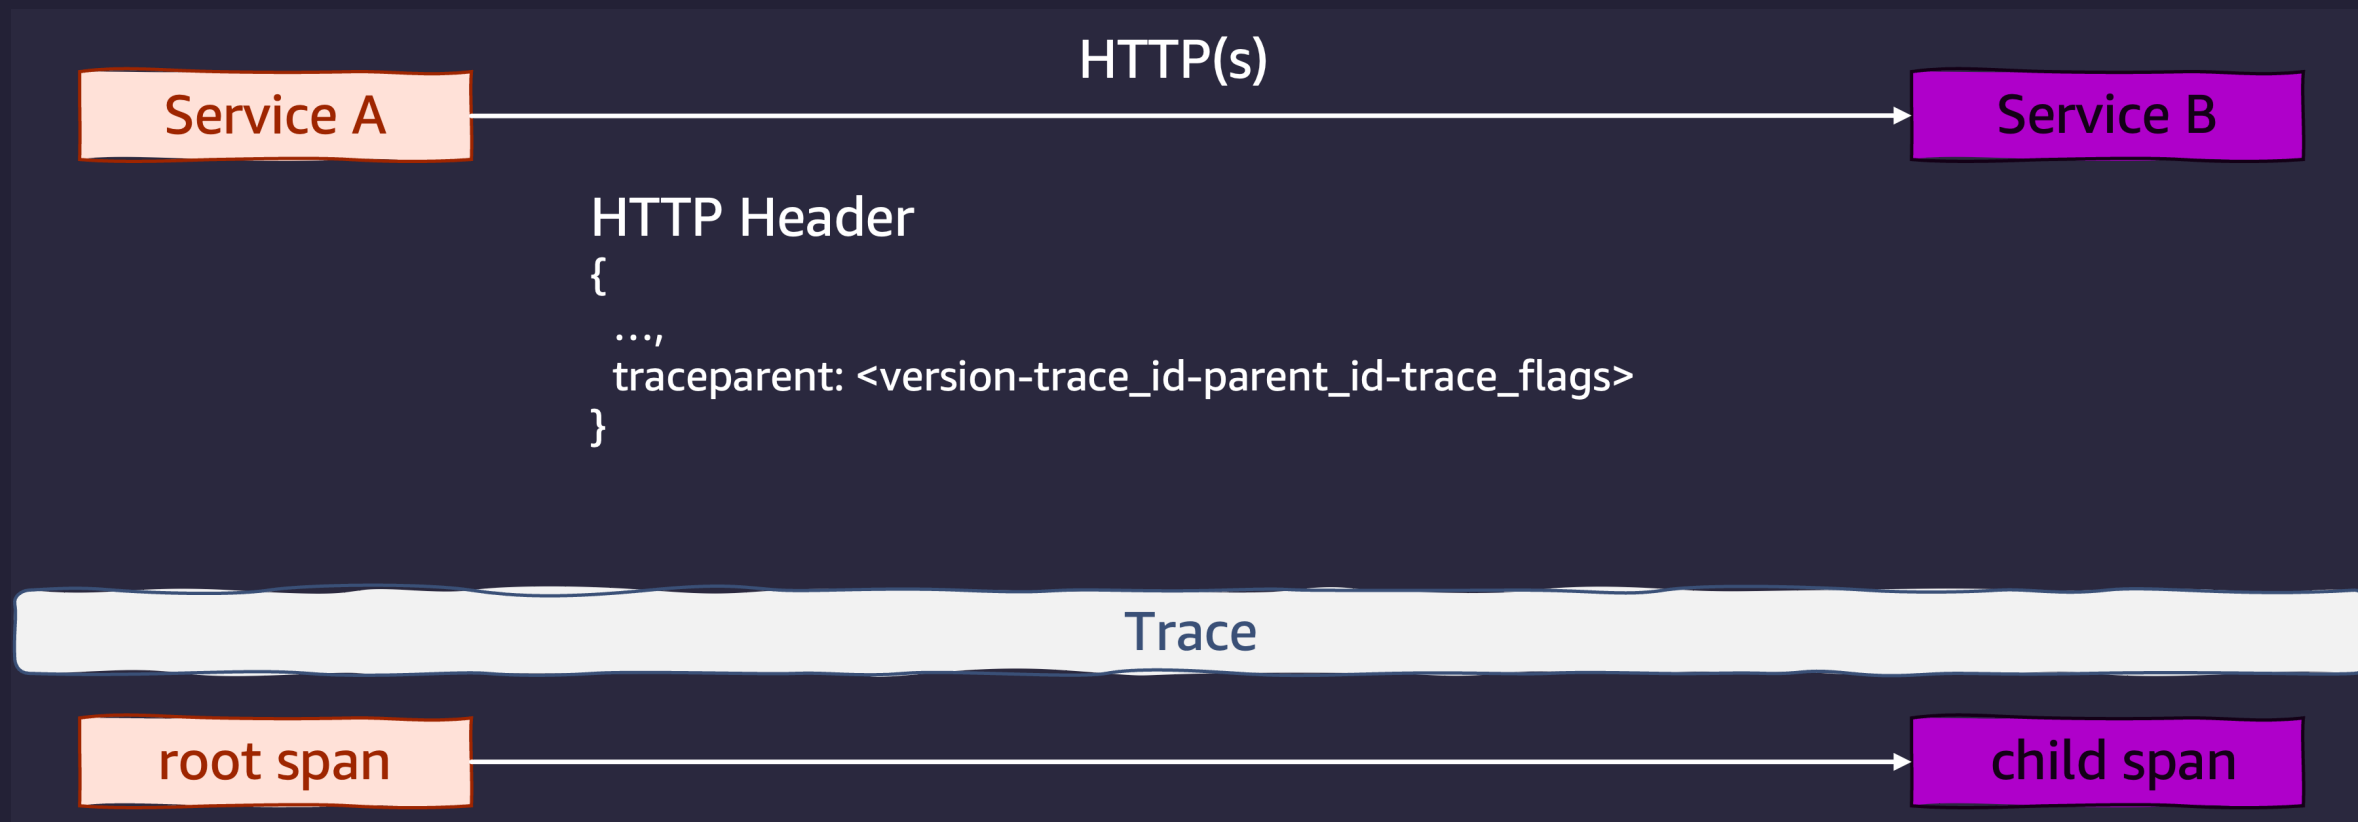
Trace



Trace (Span)



Trace - context propagation



Trace - context propagation



Trace

root span

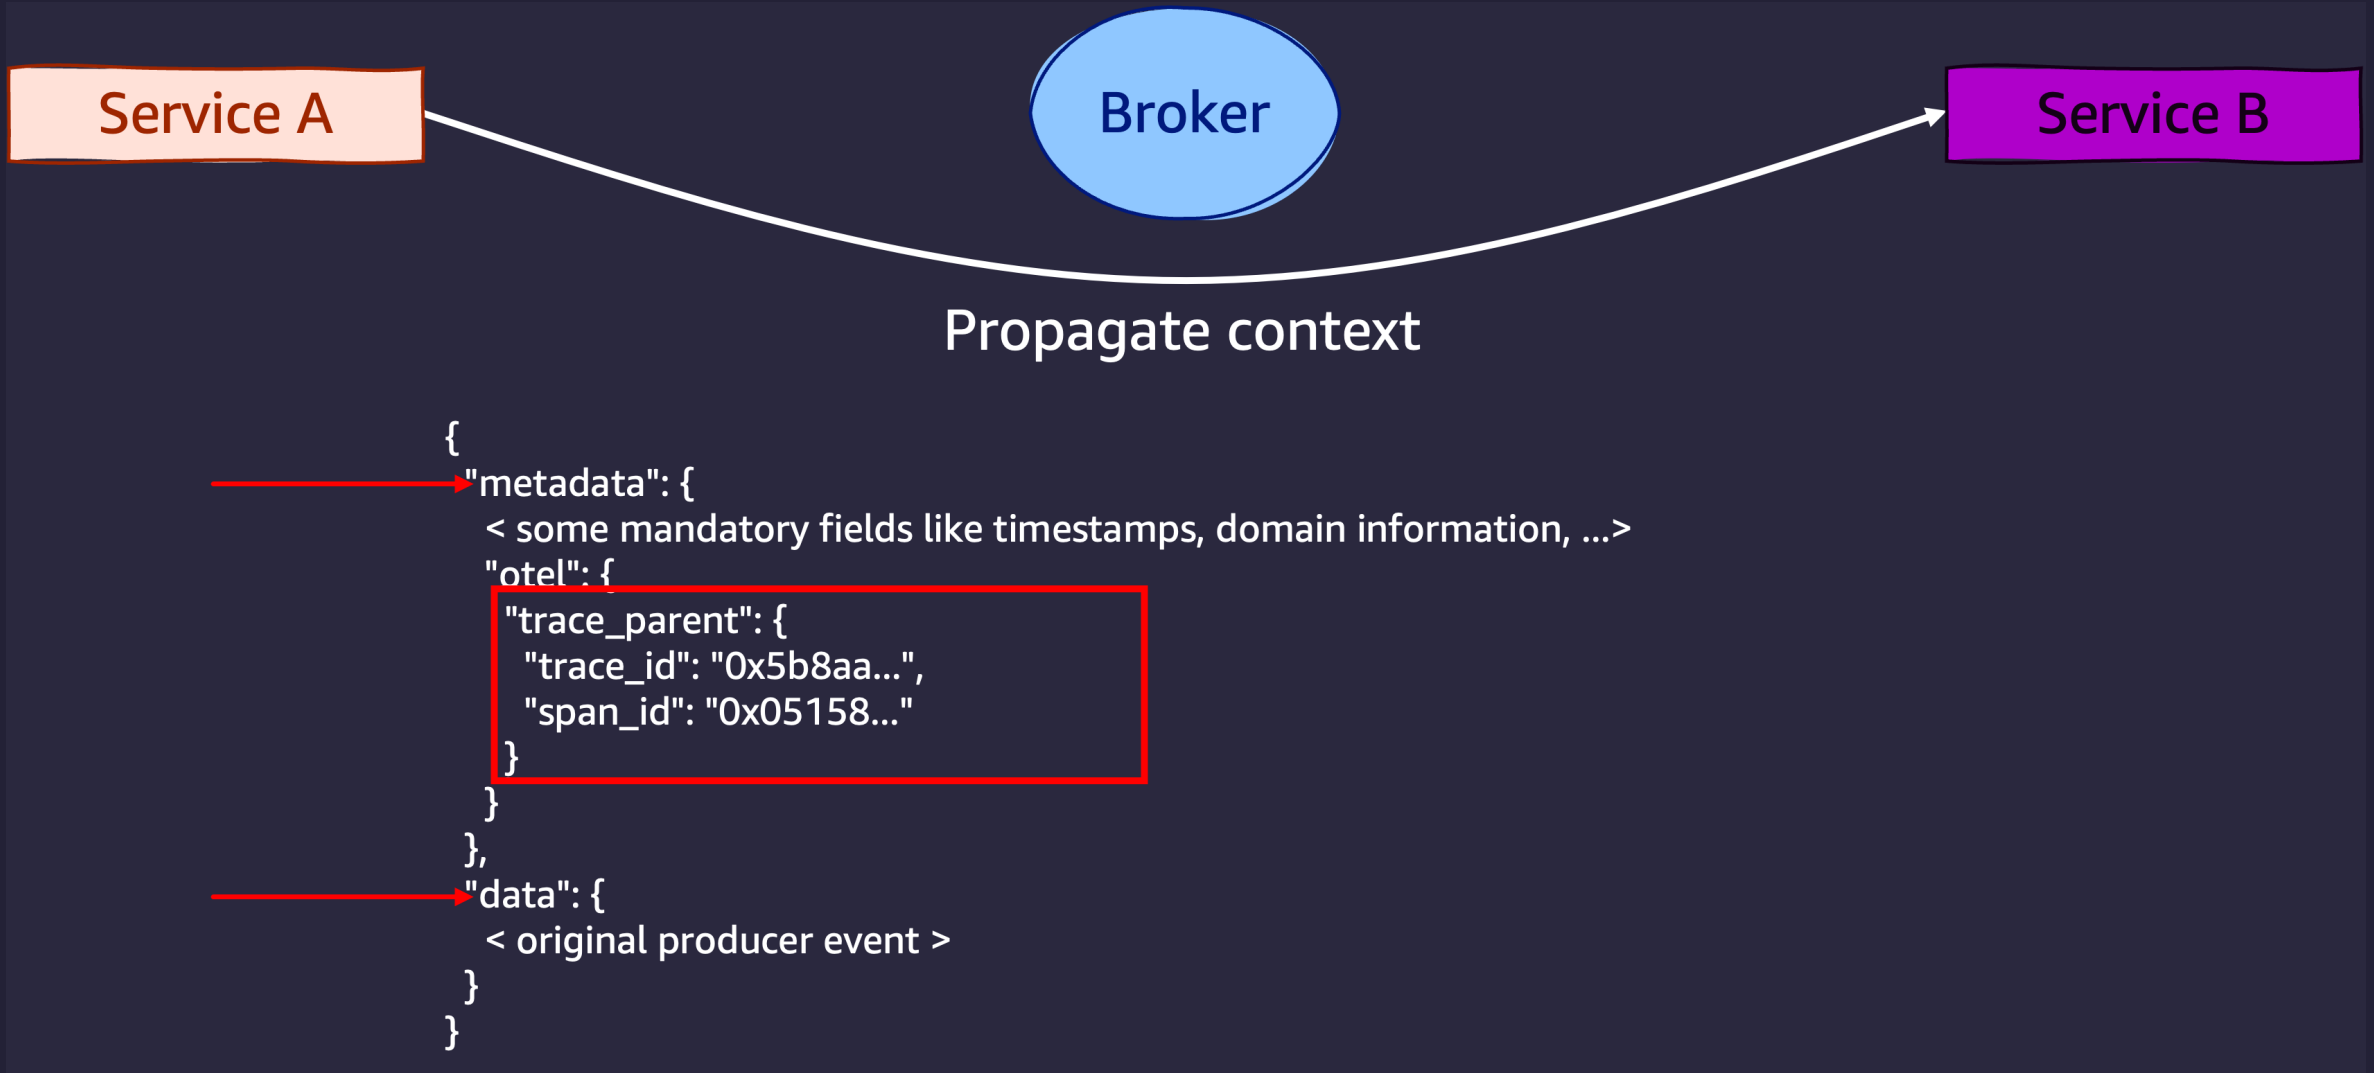
```
{
  "name": "root-span",
  "context": {
    "trace_id": "0x5b8aa5a2d2c872e8321cf37308d69df2",
    "span_id": "0x051581bf3cb55c13"
  }
  "parent_id": null,
  "start_time": "2022-04-29T18:52:58.114201Z",
  "end_time": "2022-04-29T18:52:58.114687Z"
}
```

?

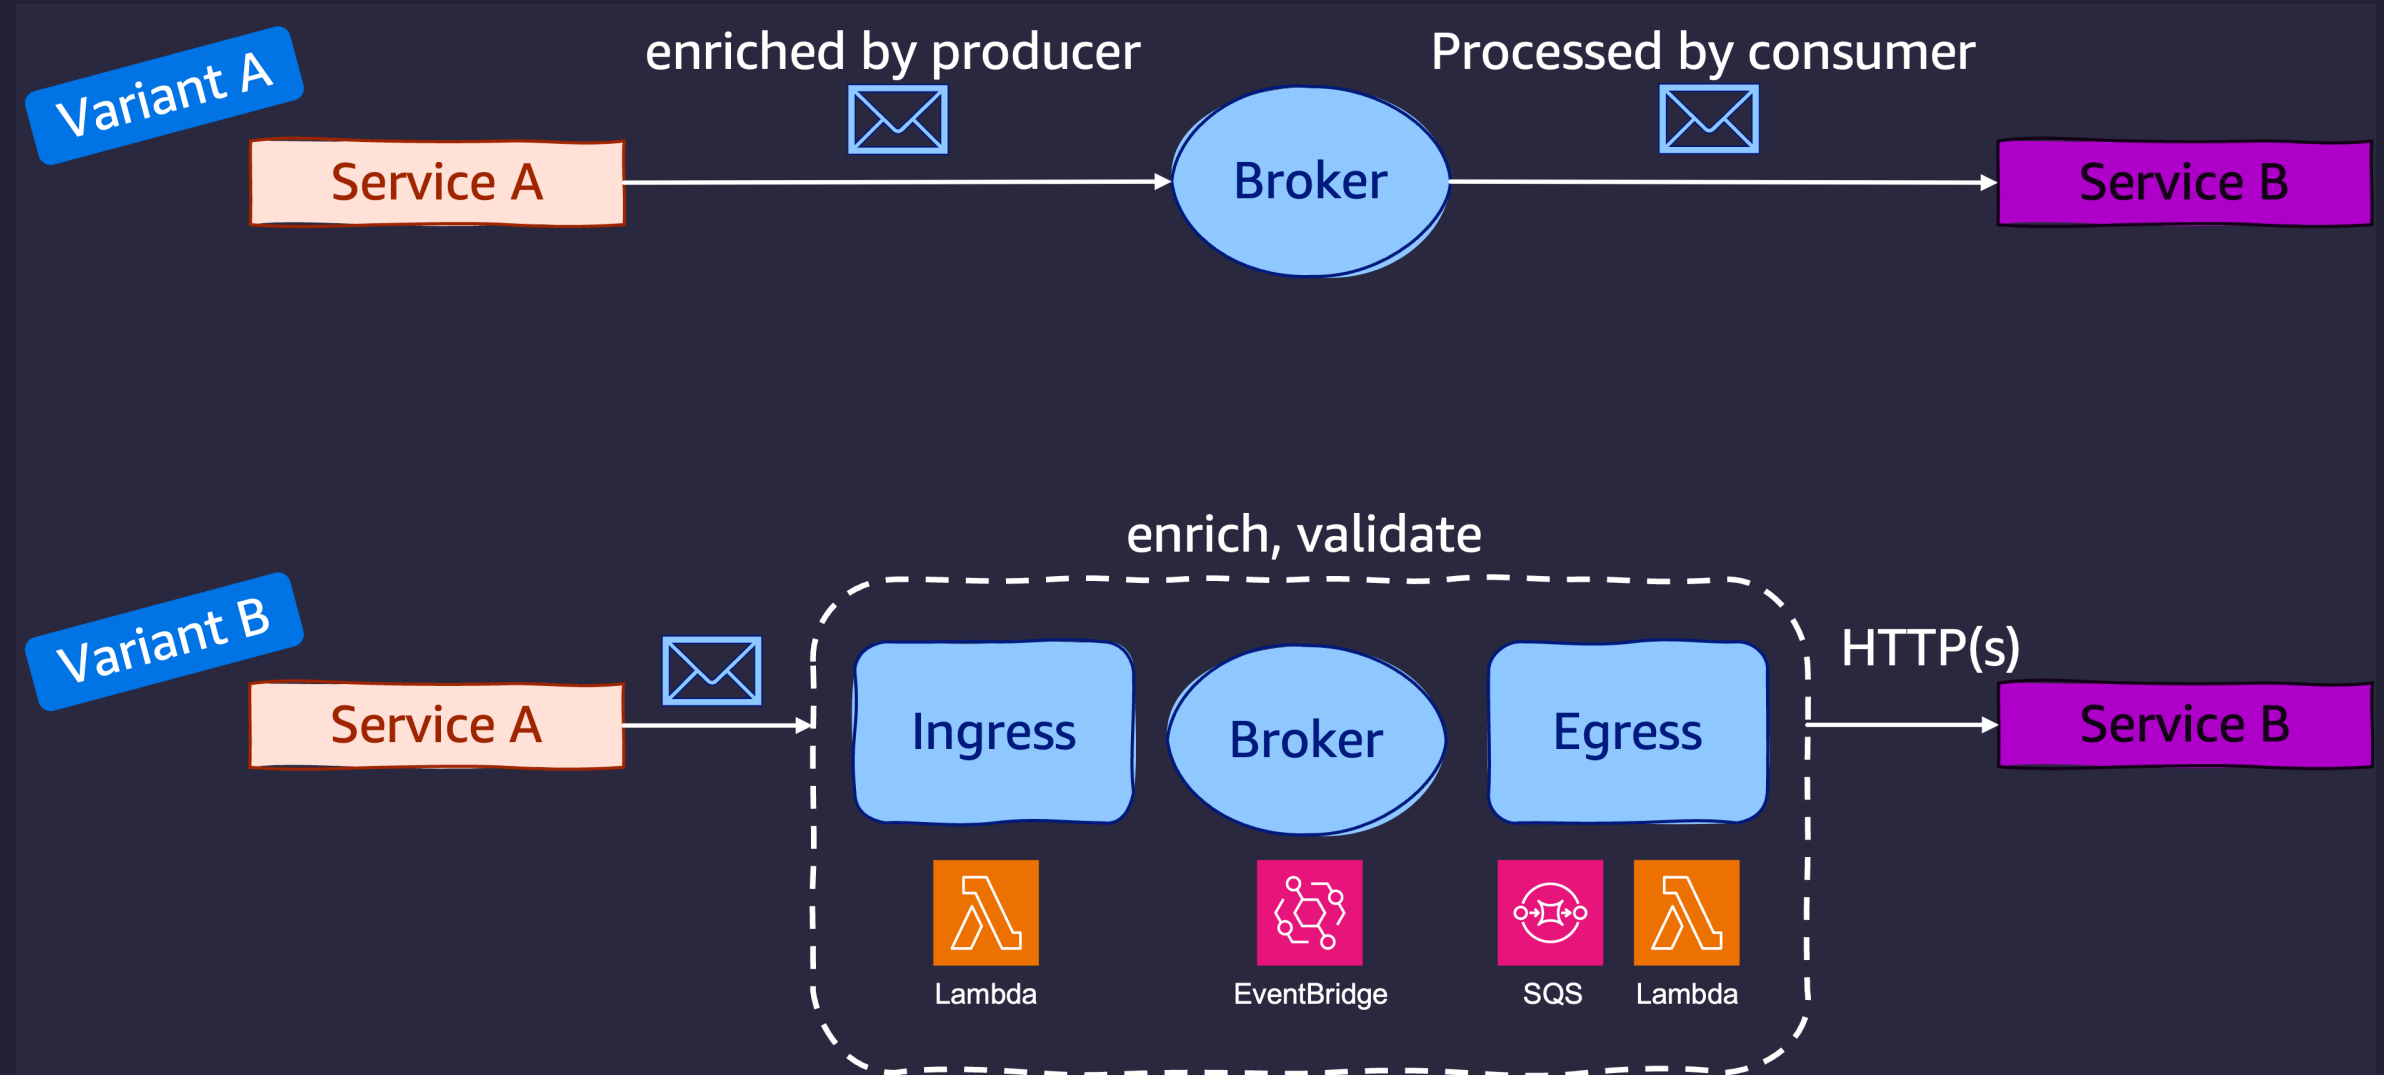
child span

```
{
  "name": "child-span",
  "context": {
    "trace_id": "0x5b8aa5a2d2c872e8321cf37308d69df2",
    "span_id": "0x5fb397be34d26b51"
  }
  "parent_id": "0x051581bf3cb55c13",
  "start_time": "2022-04-29T18:52:58.114504Z",
  "end_time": "2022-04-29T22:52:58.114561Z"
}
```

Trace - context propagation



Trace - context propagation




Span

```
{
  "name": "/v1/sys/health",
  "context": {
    "trace_id": "7bba9f33312b3dbb8b2c2c62bb7abe2d",
    "span_id": "086e83747d0e381e"
  },
  "parent_id": "",
  "start_time": "2021-10-22 16:04:01.209458162 +0000 UTC",
  "end_time": "2021-10-22 16:04:01.209514132 +0000 UTC",
  "status_code": "STATUS_CODE_OK",
  "status_message": "",
  "attributes": {
    "http.scheme": "http",
    "http.host": "10.177.2.152:26040",
  },
  "events": [
    {
      "name": "",
      "message": "OK",
      "timestamp": "2021-10-22 16:04:01.209512872 +0000 UTC"
    }
  ]
}
```

Span

170 / 521 requests | 2.5 MB / 2.8 MB transferred | 32.8 MB / 34.9 MB resources | Finish: 18.72 s | DOMContentLoaded: 1.66 s | Load: 2.54 s

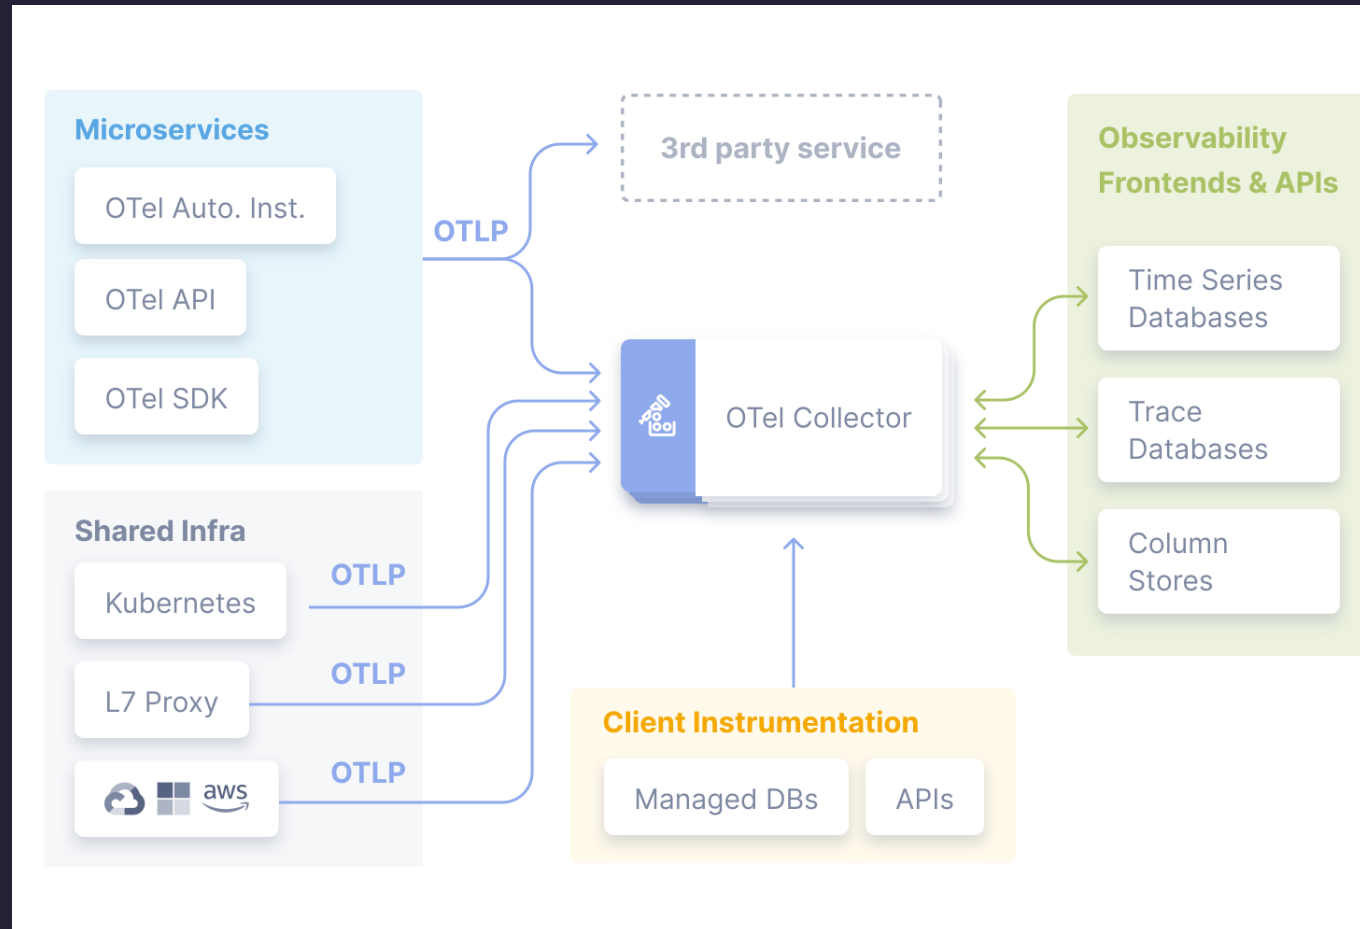
```
{
  "attributes": {
    "response_size": "2.5",
    "content_size": "2.8",
  },
  "events": [
    {
      "name": "Finish",
      "message": "OK",
      "timestamp": "2021-10-22 16:04:01.209512872 +0000 UTC"
    }
  ]
}
```



Hello, OpenTelemetry

Open source project hosted on CNCF Specifications, Implementations for instrumentation and transmissions of telemetry data (metrics, logs, traces)

1. Cross-language specifications
2. OpenTelemetry Collector (agent)
3. SDKs for each language
4. Auto Instrumentation



OpenTelemetry Instrumentation

1. Code-based solutions via official APIs and SDKs for most languages

- **API** defines data types and how to generate telemetry data.
- **SDK** defines a language-specific implementation of the API, plus configuration, data processing and exporting.

2. Zero-code solutions

- Go, .NET, PHP, Python, Java, JavaScript

Language	Traces	Metrics	Logs
C++	Stable	Stable	Stable
C#/.NET	Stable	Stable	Stable
Erlang/Elixir	Stable	Development	Development
Go	Stable	Stable	Beta
Java	Stable	Stable	Stable
JavaScript	Stable	Stable	Development
PHP	Stable	Stable	Stable
Python	Stable	Stable	Development
Ruby	Stable	Development	Development
Rust	Beta	Alpha	Alpha
Swift	Stable	Development	Development

Manually Instrumentation (Python)

```
@app.route("/server_request")
def server_request():
    with tracer.start_as_current_span(
        "server_request",
        context=extract(request.headers),
        kind=trace.SpanKind.SERVER,
        attributes=collect_request_attributes(request.environ),
    ):
        print(request.args.get("param"))
        return "served"
```


Programmatically-instrumented server (Python)

```
instrumentor = FlaskInstrumentor()

app = Flask(__name__)

instrumentor.instrument_app(app)
# instrumentor.instrument_app(app, excluded_urls="/server_request")
@app.route("/server_request")
def server_request():
    print(request.args.get("param"))
    return "served"
```

OpenTelemetry Registry

OpenTelemetry instrumentation libraries

<https://opentelemetry.io/ecosystem/registry/>

Search **NGINX**

The screenshot shows the OpenTelemetry Registry website. The header includes the OpenTelemetry logo and navigation links: Docs, Ecosystem, Status, Community, Blog, English, and a search bar. The main heading is "Registry" with the subtitle "Find libraries, plugins, integrations, and other useful tools for using and extending OpenTelemetry." Below this, a search bar shows "Search 838 entries" for the term "nginx". A description states: "The OpenTelemetry Registry allows you to search for instrumentation libraries, collector components, utilities, and other useful projects in the OpenTelemetry ecosystem. If you are a project maintainer, you can [add your project to the OpenTelemetry Registry](#)." The search results are displayed in a list. The first result is "NGINX Instrumentation" by "OpenTelemetry Authors". It is described as "NGINX OpenTelemetry module to add OpenTelemetry distributed tracing support to NGINX." The metadata on the right indicates it is a "C++ Language" "Instrumentation Component" with an "Apache 2.0 License". A "Repository" link is provided. The second result is "NGINX Native OTel Module" by "NGINX, Inc.", marked as "new" and a "first party integration". It is described as "The ngx_otel_module dynamic module enables NGINX Open Source or NGINX Plus to send telemetry data to an OTel collector." The metadata on the right indicates it is a "C++ Language" "Instrumentation Component" with an "Apache-2.0 License". Links for "Website" and "Repository" are provided at the bottom of the result card.

OpenTelemetry

Docs Ecosystem Status Community Blog English 🔍 Search this site...

Registry

Find libraries, plugins, integrations, and other useful tools for using and extending OpenTelemetry.

The OpenTelemetry Registry allows you to search for instrumentation libraries, collector components, utilities, and other useful projects in the OpenTelemetry ecosystem. If you are a project maintainer, you can [add your project to the OpenTelemetry Registry](#)

Search 838 entries nginx [Submit](#) [Reset](#) [Language](#) [Type](#)

[NGINX Instrumentation](#)

by [OpenTelemetry Authors](#)

NGINX OpenTelemetry module to add OpenTelemetry distributed tracing support to NGINX.

[Repository](#)

C++
Language

Instrumentation
Component

Apache 2.0
License

[NGINX Native OTel Module](#)

by [NGINX, Inc.](#) ★ new ♥ first party integration

The `ngx_otel_module` dynamic module enables NGINX Open Source or NGINX Plus to send telemetry data to an OTel collector.

[Website](#) [Repository](#)

C++
Language

Instrumentation
Component

Apache-2.0
License

OpenTelemetry collector



OTEL COLLECTOR

OpenTelemetry protocol (OTLP)

<https://github.com/open-telemetry/opentelemetry-proto/blob/main/docs/specification.md>

OTLP is implemented over **gRPC** and **HTTP** transports and specifies the Protocol Buffers schema used for payloads.

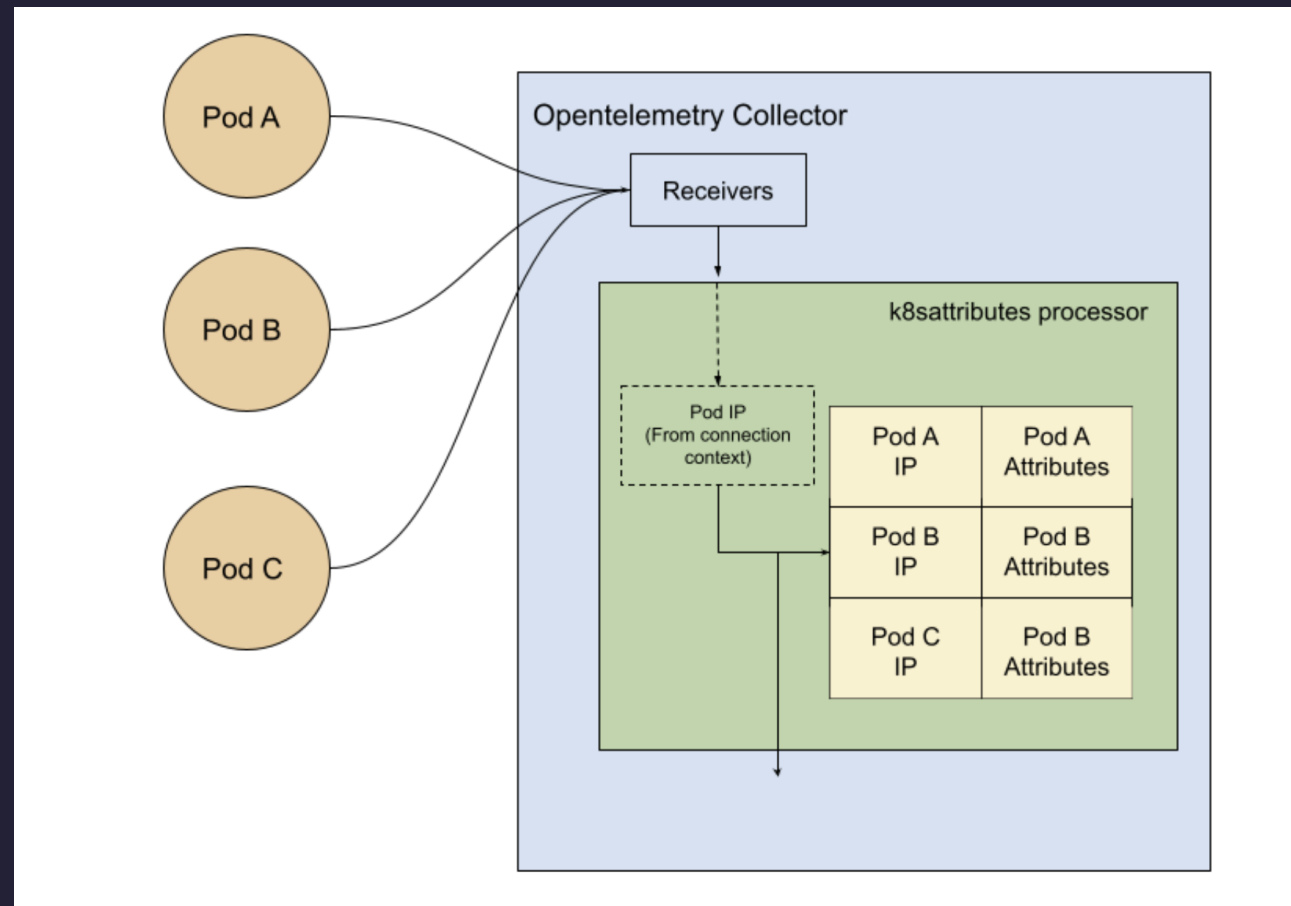
OTLP is a request/response style protocol where **clients** send requests and the **server** replies with corresponding responses.

All server components must support the following transport compression options: **none**, **gzip**

Resource Semantic Conventions

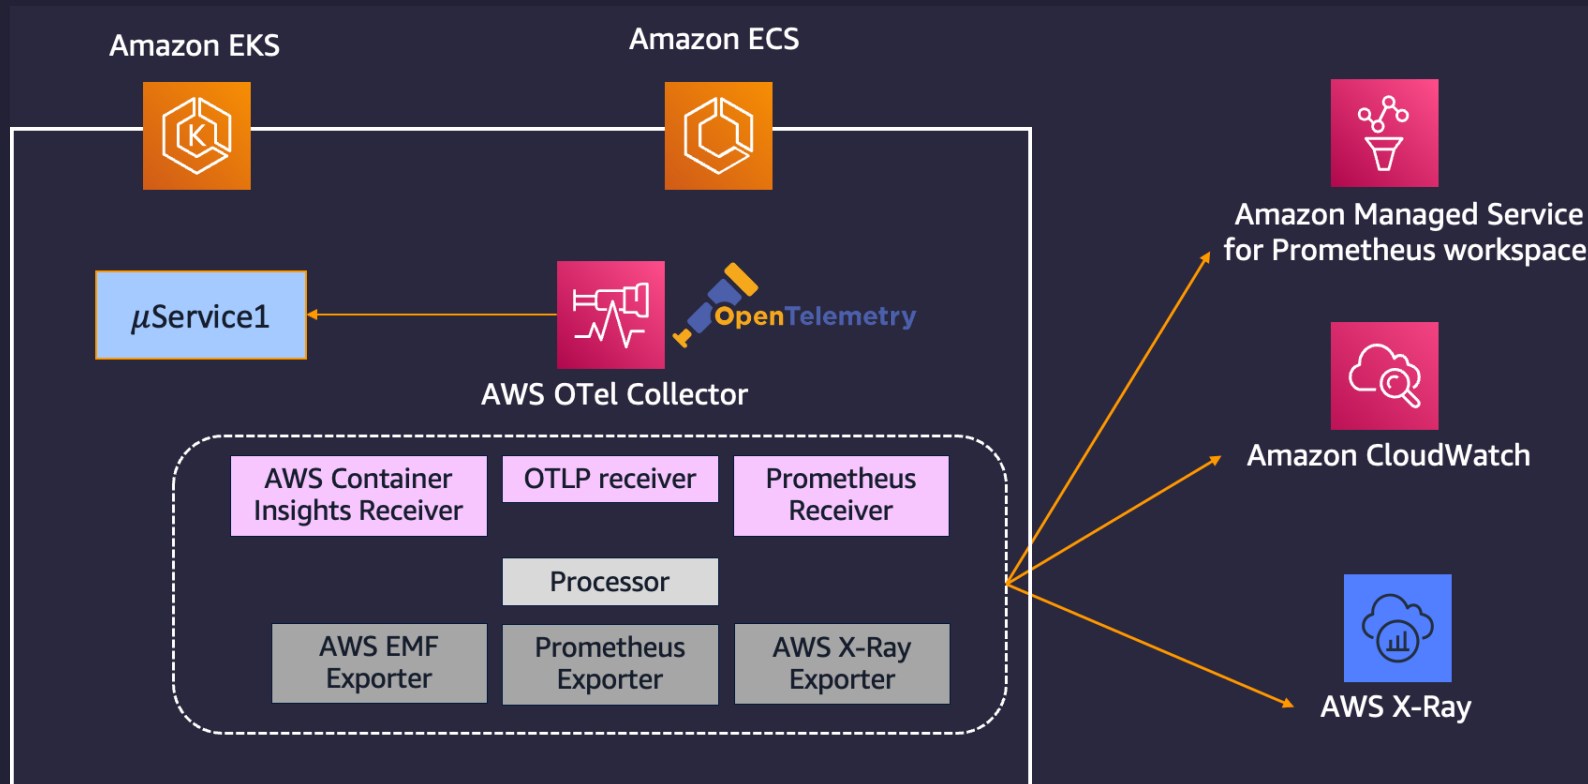
A **Resource** represents the entity producing telemetry as resource attributes.

<https://opentelemetry.io/docs/specs/semconv/>




AWS Distro for OpenTelemetry (ADOT)

- Secure, production ready, and supported by AWS OpenTelemetry distribution
- Extend functionality for ease of use on AWS



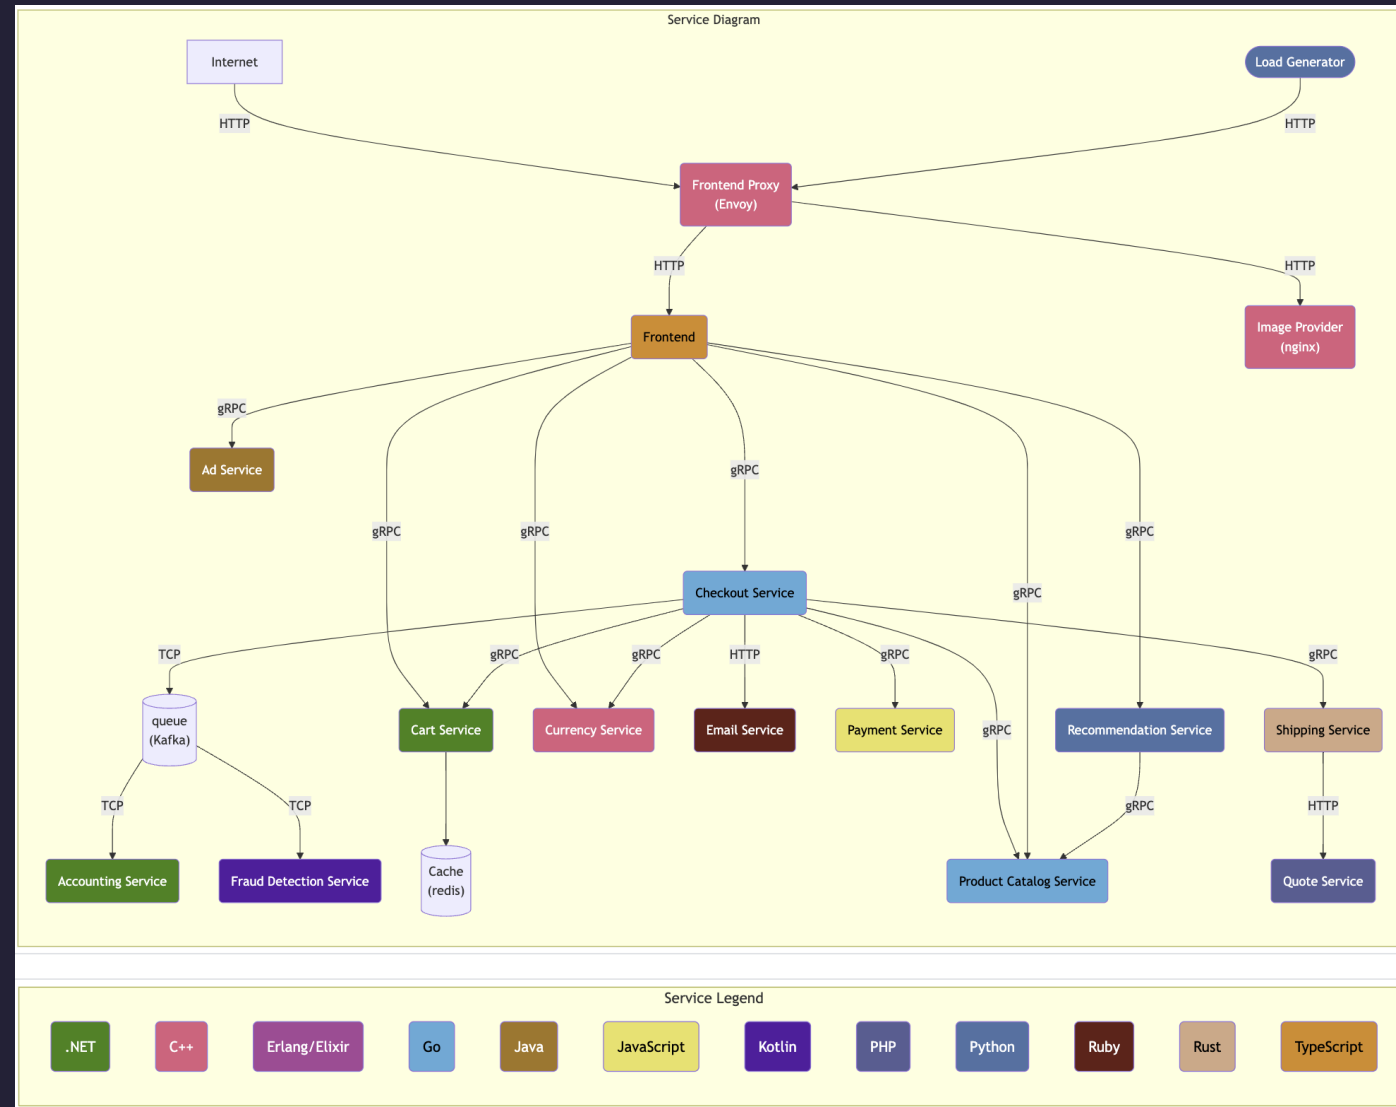
AWS Distro for OpenTelemetry (ADOT)



OpenTelemetry Demo

OpenTelemetry Demo is composed of microservices written in different programming languages that talk to each other over gRPC and HTTP; and a load generator which uses Locust to fake user traffic.

- Web store
- Grafana
- Load Generator
- Jaeger UI



New Otel Feature

Envoy and Istio

Profiling Agent

LLM Observability

OpenTelemetry Collector Antipatterns

References

- <https://opentelemetry.io/docs/>
- <https://w3c.github.io/trace-context/>
- <https://w3c.github.io/baggage/>
- <https://github.com/open-telemetry/opentelemetry-specification>
- <https://opentelemetry.io/docs/specs/semconv/>
- <https://opentelemetry.io/docs/specs/otel/protocol/>
- <https://opentelemetry.io/docs/concepts/sampling/>
- <https://opentelemetry.io/docs/demo/>
- <https://opentelemetry.io/blog/2024/>

Thank You